

Project 1 (FYS-STK4155)

Yoan Tutunarov & Vladislav Foss & Sebastian Koranda

October 2025

Abstract

This project investigates three regression methods—Ordinary Least Squares (OLS), Ridge, and LASSO—applied to polynomial fitting of the Runge function. We implement both closed-form solutions and gradient descent approaches, and compare their performance, noting small deviations that are expected due to numerical optimization. To analyze model complexity, we study the bias–variance trade-off using bootstrap resampling, which provides estimates of bias, variance, and noise. The results show that OLS overfits at higher polynomial degrees, Ridge reduces variance by shrinking coefficients, and LASSO promotes sparsity by eliminating irrelevant terms. Overall, the study illustrates how regularization and resampling techniques improve predictive performance and provide deeper insight into the generalization properties of regression models.

Acknowledgments and AI Disclaimer

We would like to thank our supervisor for valuable feedback and discussions throughout the project. This project has also made limited use of AI tools, such as ChatGPT, as a support mechanism for improving clarity, refining explanations, and troubleshooting code. The primary analysis, problem solving, and discussion were carried out independently by the project members. AI assistance was used only to enhance language in the report, verify understanding, and used for polishing figure names, ensuring that all team members fully comprehend the underlying concepts and methods.

Introduction

In this project, we study different regression methods with a particular focus on Ordinary Least Squares (OLS), Ridge regression, and Lasso regression.

These methods form the foundation of many statistical and machine learning approaches, and they provide useful insights into the trade-off between model complexity and predictive performance. As a test case, we apply these methods to the Runge function, a one-dimensional function that is well known to illustrate the challenges of polynomial interpolation, especially at higher polynomial degrees (Runge's phenomenon). By fitting polynomials of varying complexity to this function, we investigate how different regression techniques handle overfitting, variance, and bias. The project is based on weekly exercises, lecture material, and additional references. All numerical experiments are implemented in Python, with the code provided in a separate file and linked through GitHub. This report serves as documentation of our implementation, results, and analysis.

Theory, Methods and Concepts

Mean Squared Error (MSE)

Mean Squared Error (MSE) is one of the most fundamental measures of model performance in statistics and machine learning. It quantifies how well a predictive model approximates the true data by measuring the average squared difference between the actual (target) values and the predicted values [2].

Formally, for a dataset with n samples, the MSE is defined as

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2,$$

where

- y_i denotes the true (actual) value of the response variable for observation i ,
- \hat{y}_i denotes the predicted value from the model for the same observation,
- n is the total number of observations.

The term $(y_i - \hat{y}_i)$ is called the residual, which represents the error between the actual and predicted values for observation i . Squaring the residual ensures that negative and positive deviations are treated equally and emphasizes larger errors. Averaging over all n samples provides a single metric that reflects the model's overall predictive accuracy.

A lower MSE indicates that the model predictions are closer to the actual values, while a higher MSE suggests larger discrepancies. Therefore, in

regression problems, the objective is typically to minimize the MSE, ideally approaching zero.

R² score

The R^2 score (often called the coefficient of determination) is a standard metric in regression analysis that quantifies how well the variation in the independent variables explains the variation in the dependent variable. In simpler terms, R^2 measures the proportion of the variance in the target y that is captured by the regression model, relative to a baseline model that always predicts the mean of y [10].

Formally, let y_i be the actual observed values, \hat{y}_i the predicted values, and $\bar{y} = \frac{1}{n} \sum_{i=1}^n y_i$ the sample mean. Then

$$\text{SSE} = \sum_{i=1}^n (y_i - \hat{y}_i)^2, \quad \text{SST} = \sum_{i=1}^n (y_i - \bar{y})^2,$$

and

$$R^2 = 1 - \frac{\text{SSE}}{\text{SST}} = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \bar{y})^2}.$$

Here:

- SSE (“Sum of Squared Errors”) is the sum of squared deviations between predictions and actual values.
- SST (“Total Sum of Squares”) is the sum of squared deviations of actual values around their mean.
- R^2 ranges (in theory) from $-\infty$ to 1. An $R^2 = 1$ indicates perfect prediction (all points lie exactly on the regression line). An $R^2 = 0$ indicates that the model is no better than the constant mean predictor. Negative R^2 occurs when the model performs worse than simply predicting the mean.

A higher R^2 implies a better-fitting model in that it reduces residual error relative to the variance of the data itself.

Regression

Regression is a statistical method used to study the relationship between a dependent variable (outcome) and one or more independent variables (predictors). It estimates how changes in predictors are associated with changes in the outcome. In its simplest form, linear regression fits a straight line through

data to best explain this relationship. More complex forms include multiple regression (many predictors) and nonlinear regression (curved relationships). Regression is widely used for prediction, trend analysis, and decision-making in fields like economics, finance, and machine learning [1].

Ordinary Least Squares (OLS)

Ordinary Least Squares (OLS) is a method used to estimate the coefficients of a linear regression model. It finds the values of the coefficients that minimize the total squared distance between the observed data points and the values predicted by the model. Mathematically, OLS solves:

$$\min_{\beta} (Y - X\beta)'(Y - X\beta)$$

where Y is the vector of observed values, X is the matrix of predictors, and β is the vector of coefficients. OLS relies on certain conditions: the errors should be independent, have constant variance, and ideally follow a normal distribution. The method can struggle when predictors are highly correlated or when the dataset is too small compared to the number of variables. Once estimated, OLS coefficients let us assess model quality (e.g., using R^2) and make predictions [3]. In Task A, we implemented our own OLS function to compute the optimal $\hat{\theta}$. From the lecture notes in Week 35, we know that the solution to the regression problem is given as $y = X\theta$, where X is the design matrix. In this assignment we used a polynomial of degree 15, meaning that $\hat{\theta}$ is a one-dimensional array with 15 parameters.

The solution for OLS is obtained by differentiating the cost function and solving for its minimum:

$$\begin{aligned}\frac{\partial C}{\partial \theta} &= 0 = X^\top (y - X\theta), \\ X^\top y &= X^\top X\theta,\end{aligned}$$

and given that $X^\top X$ is invertible, the solution becomes

$$\hat{\theta} = (X^\top X)^{-1} X^\top y.$$

Here y is given by Runge's function $y = \frac{1}{1+25x^2}$ with added noise. The noise was generated using NumPy's `np.random.normal()`, with parameters specified in the project description.

We then constructed the design matrix X and split the data into training and test sets using scikit-learn's built-in `train_test_split()` function.

Ridge regression

Ridge Regression (also called L2 regularization) is a variation of linear regression that introduces a penalty on large coefficient values in order to reduce model variance and improve stability in the presence of multicollinearity. It adjusts the ordinary least squares estimator by adding a term λI (where $\lambda \geq 0$ is the regularization parameter and I is the identity matrix) such that the coefficient estimates are given by

$$\hat{\beta}_\lambda = (X^\top X + \lambda I)^{-1} X^\top y$$

This estimator shrinks the coefficients toward zero, but unlike some other methods, does not set any coefficients exactly to zero. The choice of λ controls the bias-variance tradeoff: a larger λ increases bias but lowers variance and may improve prediction for new data. Care must be taken: if λ is too large, the model may underfit; if too small, it may behave similarly to OLS and still overfit [4].

LASSO Regression

LASSO (Least Absolute Shrinkage and Selection Operator) regression is an extension of linear regression that introduces both regularization and variable selection. It modifies the ordinary least squares (OLS) objective by adding an L_1 penalty term, which constrains the absolute values of the model coefficients. The LASSO objective function is defined as

$$\min_{\beta} \left\{ \frac{1}{2n} \sum_{i=1}^n (y_i - \mathbf{x}_i^\top \beta)^2 + \lambda \sum_{j=1}^p |\beta_j| \right\},$$

where $\lambda \geq 0$ is a regularization parameter controlling the strength of the penalty.

The L_1 regularization term encourages sparsity in the coefficient vector β by shrinking some coefficients exactly to zero, thereby performing feature selection. This simplifies the model and helps manage the bias-variance trade-off: as λ increases, model variance decreases while bias increases. However, if λ is too large, the model may underfit and fail to capture important data patterns.

Unlike OLS and Ridge regression, LASSO does not admit a closed-form analytical solution because the L_1 norm introduces a non-differentiable point at zero. Consequently, iterative optimization algorithms such as coordinate descent or gradient descent are required to estimate the parameters [5].

Gradient Descent Variants

In machine learning, we typically have a dataset X and a model defined by parameters θ (sometimes also denoted β). The objective is to determine the parameter values that minimize a chosen cost function $C(\theta)$. Gradient descent is an iterative optimization algorithm designed to solve this problem.

The main idea is the following: if we have a function $F(x)$, with $x = (x_1, x_2, \dots, x_n)$, then the gradient $\nabla F(x)$ points in the direction of the steepest ascent of the function. By moving in the opposite direction of the gradient, the algorithm approaches a minimum of the cost function after taking enough steps. The general update rule for gradient descent is

$$\mathbf{x}_{k+1} = \mathbf{x}_k - \eta_k \nabla F(\mathbf{x}_k),$$

where the parameter η is the *learning rate*. This can be thought of as the step size taken in the chosen direction. If η is too small, convergence towards the minimum is very slow, while if η is too large, the algorithm may diverge. Choosing a suitable learning rate is therefore of great importance.

In our project, the goal is to minimize the Mean Squared Error (MSE) cost function for OLS, and MSE with an additional penalty term for Ridge regression. Minimization is performed by iteratively adjusting the model parameters θ , thereby reducing the error between the predicted values and the true outcomes [6]. We start by initializing all coefficients to zero. At each iteration, we compute the gradient of the cost function, which tells us how the error changes with respect to each parameter. The gradient indicates the slope of the loss surface: a positive gradient means the parameter should be decreased, while a negative gradient means it should be increased. Using a fixed learning rate η , we then update the coefficients according to

$$\theta \leftarrow \theta - \eta \nabla_{\theta} C(\theta).$$

For OLS, the gradient is

$$\nabla_{\theta} C_{\text{OLS}}(\theta) = \frac{2}{n} X^{\top} (X\theta - y),$$

while for Ridge it becomes

$$\nabla_{\theta} C_{\text{Ridge}}(\theta) = \frac{2}{n} X^{\top} (X\theta - y) + 2\lambda\theta,$$

where λ is the regularization parameter [7].

Stochastic Gradient Descent (SGD)

In the most basic version of the gradient descent, we usually include the whole dataset to compute the the gradient descent at each stop. This is in contrast to **SGD**, where we would calculate the gradient on a subset of data called mini-batches. If our whole dataset is \mathbf{n} datapoints, and we want to have a mini-batch of size \mathbf{M} , then the number of batches would be $\frac{\mathbf{n}}{\mathbf{M}}$.

Compared to the full batch (whole dataset) gradient descent, SGD typically will converge faster per iteration, but will however experience larger variance due to noise and the random nature of the method. A trade off with the SGD is the size of each batch: larger batches will be computationally more expensive but less noisy, and smaller batches will be very cost efficient, however will have a lot more noise in the trajectory they are calculating [9].

The gradient step of the SGD is calculated as:

$$\theta_{t+1} = \theta_t - \eta_t \nabla C_{B_k}(\theta_t),$$

Here k represents the randomly chosen mini-batch from the dataset at each iteration. One complete pass over all mini-batches is referred to as an *epoch* [9]. Calculating the gradient on subsets of the data reduces computational cost compared to full-batch gradient descent, and the inherent randomness can also help the algorithm escape shallow local minima and saddle points.

Since the learning rate plays a critical role in the performance of SGD, a variety of adaptive learning rate algorithms have been developed. These methods adjust the step size dynamically based on information from previous updates. In this project, we focus on four variants: **Momentum**, **AdaGrad**, **RMSProp**, and **Adam**, each of which introduces specific improvements to the basic GD and SGD functions.

Momentum based GD

SGD generally today is implemented with momentum. The general idea behind momentum based GD / SGD is that if the landscape we are traversing with our gradient is a canyon, the vanilla gradients will have large oscillations instead of gliding towards the minima. The implementation of momentum build on the idea that if the direction of the gradient descent aligns with the previous point, then we gain "momentum", and if not we dampen it. This way we dampen the oscillations that we might experience. This is

mathematically described as:

$$\begin{aligned}\theta_{t+1} &= \theta_t - v_t, \text{ where} \\ v_t &= \gamma v_{t-1} + \eta_t \nabla E(\theta_t)\end{aligned}$$

The function $E(\theta)$ represents our function we are trying to find the minimum of [9]. Here we can see that we will get a smoothed path where we accelerate if the gradient is still pointing in the same direction. We also introduce a new parameter γ , which represents the momentum, and we usually choose between $0 \leq \gamma \leq 1$ [9].

Adagrad

The adaptive learn rate of Adagrad drops the parameter γ represented in momentum completely and has a different approach. For this algorithm, we first define a vector:

$$r_t = r_{t-1} + (\nabla C(\theta))^2$$

Calculating the AdaGrad at step t will then become

$$\theta_{t+1} = \theta_t - \eta H_t^{-1/2} \nabla C(\theta)$$

, where we have that the matrix H_t is defined as $H_t = \text{diag}(r_t)$, where the entries in the matrix will diagonally be entries of r_t over time. The update of θ will be:

$$\theta_{t+1,j} = \theta_{t,j} - \frac{\eta}{\sqrt{\epsilon + r_{t,j}}} \nabla C(\theta_{t,j})$$

Here we introduce the parameter ϵ as a small constant to avoid division by zero if $r_{t,j}$ is 0. From the mathematical expressions, we can see that there we have removed the momentum (γ) as implemented in momentum gradient descent described in a section above. We see that if a gradient is very large the step size we take ($\frac{\eta}{\sqrt{\epsilon + r_{t,j}}}$) will automatically become very small, meaning we have less large jumps, while if the gradient is very small the step size becomes larger and speeds up [9].

Unfortunately, a key limitation of AdaGrad is that the accumulated squared gradients in H_t grow without bound. This causes the algorithm to take very large steps in the beginning, but as the vector $r_{t,j}$ grows, the effective learning rate

$$-\frac{\eta}{\sqrt{r_{t,j} + \epsilon}}$$

becomes smaller and smaller [9]. As a result, the algorithm slows down excessively, and the step sizes eventually become infinitesimally small, effectively stopping further learning. This limitation is the main motivation behind the development of RMSProp and Adam, which modify AdaGrad to maintain a more stable learning rate [9].

RMSProp

Subsequently, the root-mean-square propagation (RMSProp) algorithm was developed to address AdaGrad’s limitation of diminishing learning rates that eventually lead to extremely small step sizes. The key idea is to use an exponentially decaying moving average of past squared gradients rather than accumulating all of them. This acts as a weighted average that gives more importance to recent gradients than to older ones, preventing the learning rate from decaying towards zero over many iterations [8, 9].

The exponentially decaying average is defined as

$$v_t = \rho v_{t-1} + (1 - \rho)(\nabla C(\theta_t))^2,$$

where ρ is the decay rate, typically chosen in the range $0.9 \leq \rho \leq 0.99$.

The update rule for the parameters then becomes

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{v_t} + \epsilon} \nabla C(\theta_t),$$

where ϵ is a small constant included to avoid division by zero.

From these expressions, we see that recent gradients have greater influence than earlier ones, ensuring that the effective learning rate remains stable. This modification makes RMSProp more robust than AdaGrad and it has become a widely used optimizer in practice, as well as an important component of the Adam algorithm [8, 9].

Adam

The adaptive moment estimation (Adam) algorithm combines the ideas of RMSProp, with its adaptive learning rate, and momentum, which smooths the optimization path. Adam maintains running averages of both the first and the second moments of the gradients [9].

The first moment is defined as

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) \nabla C(\theta_t),$$

and the second moment as

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2)(\nabla C(\theta_t))^2,$$

where the parameters are typically chosen as $\beta_1 = 0.9$ and $\beta_2 = 0.999$, with initial values $m_0 = 0$ and $v_0 = 0$ [9].

When we first calculate the first and second moments (m_t, v_t) , both are biased towards zero in the early iterations. This is because the running averages favor the most recent gradients, and thus require time to build up from zero. To correct for this bias, Adam introduces normalized versions of the moments. The bias-corrected terms are given as [9]:

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t}, \quad \hat{v}_t = \frac{v_t}{1 - \beta_2^t}.$$

The parameter update rule then becomes

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{\hat{v}_t} + \epsilon} \hat{m}_t,$$

where η is the learning rate and ϵ is a small constant to avoid division by zero.

From these expressions we see that in the first few iterations, when t is small, the denominators $(1 - \beta_i^t)$ are close to zero, leading to significantly larger values of \hat{m}_t and \hat{v}_t compared to the uncorrected terms. This correction is critical for stability in the early iterations of the algorithm [9]. As the number of iterations increases, the bias-corrected estimates converge towards the uncorrected values.

By combining the exponential moving averages from RMSProp with the smoothing effect of momentum, Adam achieves fast convergence with relatively little tuning. It is also robust when dealing with sparse gradients, making it one of the most widely used optimization algorithms in modern machine learning [9].

Bias-Variance Trade-Off

An important concept in machine learning and statistical analysis is the bias-variance trade-off. We can decompose the expected prediction error of a model into three components: bias, variance, and irreducible error (noise). This means that the expected squared error between the true values and the model's predictions can be expressed as the sum of these three components.

Knowing this, and the definition of mean squared error (MSE), we have:

$$\mathbb{E}[(\mathbf{y} - \tilde{\mathbf{y}})^2] = \text{Bias}[\tilde{\mathbf{y}}] + \text{var}[\tilde{\mathbf{y}}] + \sigma^2,$$

where

$$\text{Bias}[\tilde{\mathbf{y}}] = \mathbb{E}[(\mathbf{y} - \mathbb{E}[\tilde{\mathbf{y}}])^2],$$

and

$$\text{var}[\tilde{y}] = \mathbb{E}[(\tilde{\mathbf{y}} - \mathbb{E}[\tilde{\mathbf{y}}])^2] = \frac{1}{n} \sum_i (\tilde{y}_i - \mathbb{E}[\tilde{y}])^2.$$

We can derive this decomposition as follows. Assume the data is generated according to the model:

$$y_i = f(x_i) + \varepsilon_i, \quad \mathbb{E}[\varepsilon_i] = 0, \quad \text{Var}(\varepsilon_i) = \sigma^2,$$

and let \tilde{y}_i be the predictions of the model

$$\mathbb{E}[(y_i - \tilde{y}_i)^2] = \mathbb{E}[(f(x_i) + \varepsilon_i - \tilde{y}_i)^2].$$

Add and subtract $\mathbb{E}[\tilde{y}_i]$:

$$\begin{aligned} \mathbb{E}[(y_i - \tilde{y}_i)^2] &= \mathbb{E}[(f(x_i) - \mathbb{E}[\tilde{y}_i] + \mathbb{E}[\tilde{y}_i] - \tilde{y}_i + \varepsilon_i)^2] \\ &= \underbrace{(f(x_i) - \mathbb{E}[\tilde{y}_i])^2}_{\text{bias}^2} + \underbrace{\mathbb{E}[(\tilde{y}_i - \mathbb{E}[\tilde{y}_i])^2]}_{\text{Var}(\tilde{y}_i)} + \underbrace{\mathbb{E}[\varepsilon_i^2]}_{\sigma^2} \\ &\quad + \underbrace{2(f(x_i) - \mathbb{E}[\tilde{y}_i])\mathbb{E}[\tilde{y}_i - \mathbb{E}[\tilde{y}_i]]}_{=0} + \underbrace{2\mathbb{E}[(\tilde{y}_i - \mathbb{E}[\tilde{y}_i])\varepsilon_i]}_{=0}. \end{aligned}$$

The cross-terms vanish because $\mathbb{E}[\tilde{y}_i - \mathbb{E}[\tilde{y}_i]] = 0$ and, at test time, \tilde{y}_i is independent of ε_i with $\mathbb{E}[\varepsilon_i] = 0$. Therefore,

$$\boxed{\mathbb{E}[(y_i - \tilde{y}_i)^2] = (f(x_i) - \mathbb{E}[\tilde{y}_i])^2 + \text{Var}(\tilde{y}_i) + \sigma^2}$$

Averaging over all points $i = 1, \dots, n$ yields the dataset-level decomposition:

$$\boxed{\frac{1}{n} \sum_{i=1}^n \mathbb{E}[(y_i - \tilde{y}_i)^2] = \frac{1}{n} \sum_{i=1}^n (f(x_i) - \mathbb{E}[\tilde{y}_i])^2 + \frac{1}{n} \sum_{i=1}^n \text{Var}(\tilde{y}_i) + \sigma^2}$$

Implementation

Measurements for error

Mean Squared Error (MSE)

For calculating the Mean Squared Error (MSE), we make use of the scikit-learn library's built-in function `mean_squared_error`. This ensures both efficiency in our implementations

R^2 score

The R^2 score is implemented from its theoretical definition. The numerator and denominator are calculated separately, and then combined into the final return, as shown in Code 2. This metric will be used to evaluate the performance of our models later.

Scaling

On the topic of scaling, we usually scale our data before performing regression. This ensures that the different features θ are on comparable scales, so that no single feature dominates or is weighted less simply because of the magnitude of the feature. As a general practice in our implementations (see Listings 6, 1, and 3), we include an explicit scaling step.

The standard way we scale the data is to call the `StandardScaler` function from the Scikit-learn library [11], and create it as an instance called *scaler*. We then scale the training data and apply the same fitted scaler to the test data to avoid data contamination between the train and test split.

Regression Models

OLS for the Runge function

We begin by generating a dataset of 100 equally spaced points $x \in [-1, 1]$. For each point, the response value y is calculated from the Runge function, with added random noise to simulate measurement errors. To prepare the data for regression, we implemented the function `polynomial_features`. This function constructs the design matrix for a chosen polynomial degree p . Each row in the design matrix contains the powers of x , i.e. $[1, x, x^2, \dots, x^p]$, with the option of including or excluding an intercept term. The regression parameters are estimated with our function `OLS_parameters`, which implements the ordinary least squares (OLS) solution. Model performance is measured using two standard metrics: the mean squared error (MSE) and the coefficient of determination (R^2), both coded directly from their mathematical definitions (`R2_score`). In the main script, the design matrix is created and the dataset is split into training (80%) and test (20%) sets. The polynomial degree is varied from 1 to 15. For each degree, the following steps are carried out:

1. Construct the design matrix with `polynomial_features`.
2. Split the dataset into training and test sets.
3. Compute the OLS coefficients from the training data.

4. Use the coefficients to predict values on the test set.
5. Compute MSE and R^2 for the predictions.
6. Store the values for later comparison.

Finally, we create an array of polynomial degrees and plot both MSE and R^2 as functions of the polynomial degree. These plots provide a visual overview of how model complexity affects prediction accuracy. This implementation directly reflects the theoretical framework from the lectures: polynomial features control model complexity, the OLS solution provides a baseline estimator, and MSE and R^2 serve as performance measures. Gradient-based optimization methods (introduced later in the course) could be applied as an alternative, but here we focus on the analytical OLS approach.

Ridge

Based on the theoretical definition, we implement Ridge regression as a function `ridge_regression(X, y, λ)`. The parameters are estimated by

$$\theta = (X^\top X + \lambda I)^{-1} X^\top y.$$

To be able to study the effect of the regularization parameter λ , we selected multiple values of λ using `train_test_split` to divide data into training and test sets. For each λ , the model was fitted on training data and evaluated on the test set using mean squared error (MSE) and R^2 . The implementation is provided in Code Listing 3. We then subsequently plot the **MSE** values along with the R^2 values with the different values for λ .

We also implemented a way to gauge the difference between the **MSE** and R^2 values between the Ridge and OLS implementations, with differing λ values. The way this was done was by computing the mean of the **MSE** values for each of the λ values, and then plotting them against each other, as we can see in the code 4

Gradient Descent Variants

After running the algorithm for a sufficient number of iterations, the coefficients converge close to the closed-form solutions we derived in Parts A and B. Comparing the two approaches shows that gradient descent gives nearly the same coefficients, with small numerical differences due to the finite number of iterations and the choice of learning rate. Ridge regression, as expected, shrinks the coefficients compared to OLS.

A key observation from experimenting with the learning rate η is that it strongly influences convergence: if η is too small, the algorithm converges very slowly; if it is too large, the updates overshoot and the algorithm can diverge. Choosing a moderate value allows stable and efficient convergence. [8].

Momentum GD

We implement the momentum gradient descent as a standalone function to allow flexibility and reuse in other parts of the project. The algorithm is initialized with $\theta = 0$, and a variable **change** is introduced to store the previous update direction. At each iteration, the gradient is computed depending on whether we solve the OLS, Ridge, or Lasso case. For our momentum, we select a parameter of our choosing. We also track the mean squared error (MSE) at each iteration to monitor convergence, as shown in code listing 6.

AdaGrad

We implement the Adagrad also as a standalone function to allow flexibility and the ability to reuse in other parts of the project. We initialize $\theta = 0$ and also $\mathbf{r} = 0$. For gradient calculation we can easily choose OLS or Ridge in our function. Then we square the gradient and add it to the previous value of \mathbf{r} . We calculate the square of the gradient, the adaptive step and the **MSE** value, and return the θ , **MSE**. The code for the Adagrad implementation is seen in 6.

RMSProp

RMSProp was also implemented as a standalone function as well, allowing for reusability and flexibility. We start by initializing the parameters $\theta, \mathbf{v} = 0$, followed by setting up a matrix for storing our **MSE** values and calculating the decaying average of gradients. We then calculate the parameter \mathbf{v} , and update the θ value. The code for the RMSProp implementation can be seen in 6.

ADAM

We also implement ADAM as a standalone function, where we first initialize the parameters $\theta, \mathbf{v}, \mathbf{m} = 0$. We then loop over the iterations and calculate the 1st and 2nd moment, but also creating parameters $\hat{\mathbf{m}}, \hat{\mathbf{v}}$ for the bias

corrections, and using these in the final update for θ . The code for the ADAM implementation can be seen in 6.

Stochastic Gradient Descent (SGD)

We implemented a Stochastic gradient descent (SGD) as a standalone function with different adaptive learning rates that we tested, along with **OLS** and **Ridge** regression methods. The implementation follows the principles laid out in the section about the SGD method, further above.

At each epoch, the data is shuffled randomly to make sure the batches are not correlated between the updates. Then we have the minibatches being created with the chosen length of the batches. For every minibatch, the gradient is computed based on the chosen style, and further, a learn rate optimizer is being used. The ones that are in question are the **Momentum**, **AdaGrad**, **RMSProp** and **ADAM**. The function then inside of these optimizers calculated the θ update, and also calculates the **MSE** values for the minibatch. This way we can check convergence later in plotting.

The learn rate was set to $\eta = 0.03$, with a momentum factor of 0.3, as well as the λ being able to be specified. The reasons for setting these parameters was that they during testing provided most stable plots and convergences. The full implementation of the SGD is shown in code listing 8

Resampling

Bootstrap setup

To perform bootstrap resampling, we call the `resample` utility function of scikit-learn iteratively on the training data for the number of bootstraps given as argument. This function samples the dataset, with replacement, meaning it returns a dataset of the same size, consisting of elements randomly selected from the original dataset while allowing duplicates. Essentially, we get variations of the original data with some points repeated and others omitted. Repeating this process multiple times generates a distribution of datasets that reflect the variability inherent in the original data.

For each iteration, we fit the model using the resampled data and perform a prediction using each resulting fit. After completing all bootstrap iterations, we compute the mean and standard deviation of the predictions across all bootstrap samples. This provides an estimate of the uncertainty in our model's predictions, as the variability across the bootstrap samples reflects

the sensitivity of the model to changes in the training data.

Bootstrap resampling is implemented as a method of the OLS class in Appendix B 9

Cross-validation setup

For cross-validation, we use the `KFold` function from scikit-learn to split the data into k folds. For each fold, we fit the model on $k - 1$ folds and validate it on the remaining fold. We repeat this process for all folds and compute the average performance metrics (MSE and R^2) across all folds to assess the model's generalization ability.

K-folds cross validation is also implemented for the OLS class 9

Results

Part A: Ordinary Least Square (OLS) for the Runge function

We fitted the noisy Runge function $f(x) = 1/(1 + 25x^2)$ using OLS with polynomial features up to degree 15. We tested both the scaled and unscaled versions of the feature matrix to evaluate the effect of scaling the data. The results are from the code listed in 1. The results can be seen in figure 1.

Part B: Adding Ridge regression for the Runge function

We implemented the **Ridge** regression model and fitted it to the noisy Runge function $f(x) = 1/(1 + 25x^2)$ using different values of the regularization parameter λ . The tested values were

$$\lambda = [0.0001, 0.01, 0.1, 0.5].$$

The data were scaled prior to fitting, and the full implementation and plotting code are provided in Code Listing 3. The resulting fits are shown in Figure 2.

We also plotted the MSE values of OLS vs Ridge by taking the average of the values for the dataset, and then subtracting the differences between them for each of the λ parameters. The results can be seen in figure 3, and the code for implementation can be found in code listing 4.

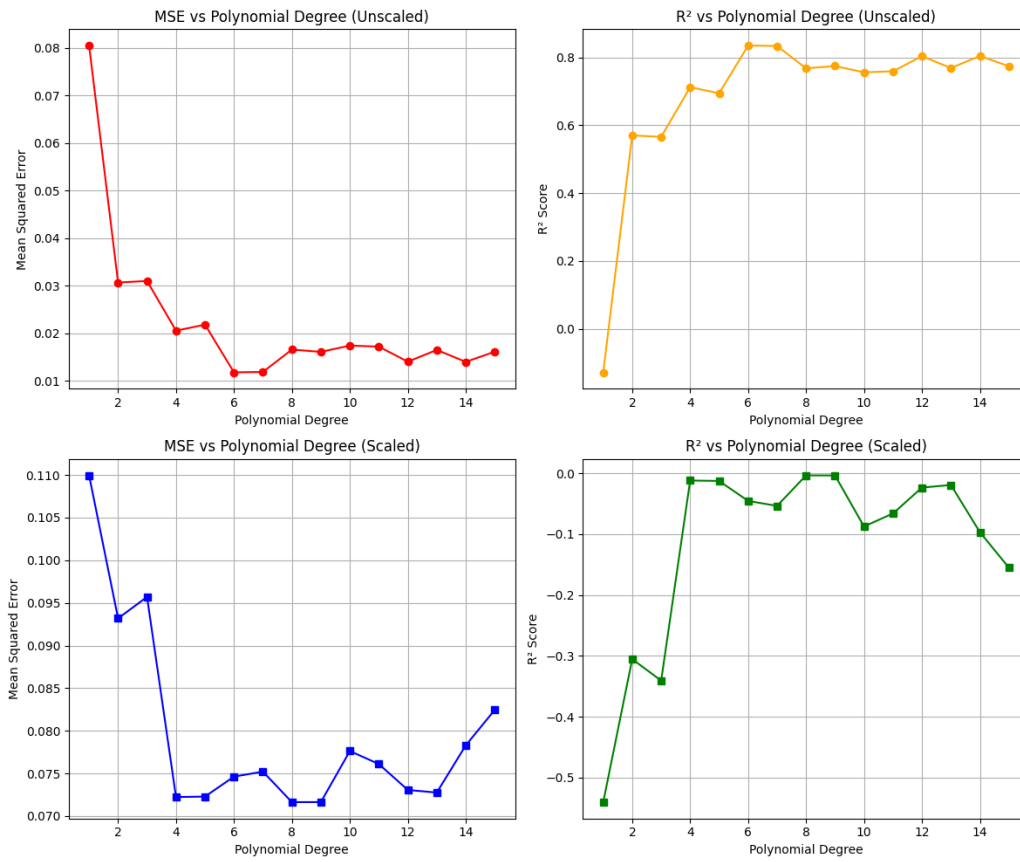


Figure 1: OLS and Ridge regressions. Scaled and unscaled versions

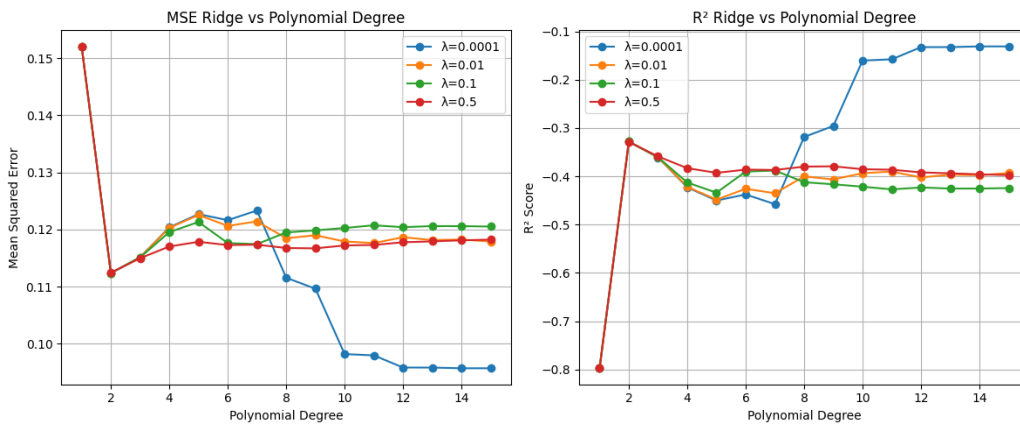


Figure 2: Ridge Analysis for different polynomial degrees




Figure 3: OLS vs Ridge comparisson in MSE

Part C: Writing our own Gradient descent code




Figure 4: GD Convergence

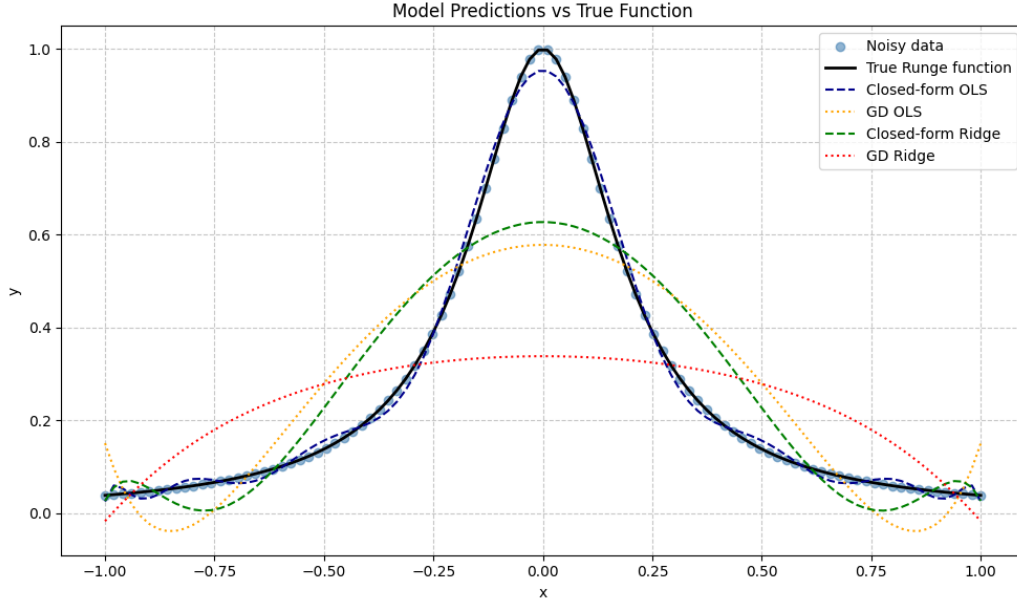


Figure 5: GD: Model prediction vs True function

Part D: Including momentum and more advanced ways to update the learning the rate

For this assignment, we applied the different adaptive gradient descent optimizers (**Momentum**, **AdaGrad**, **RMSPProp**, and **Adam**) to both **OLS** and **Ridge** regression on the noisy Runge function $f(x) = 1/(1 + 25x^2)$ using polynomial features of degree 6. In testing this was proven to give the best fit. The models were trained on the training data (80/20 split), and the fitted parameters were then applied to the entire dataset to visualize the final approximation to the true function.

The learning rate was set to $\eta = 0.03$ with a momentum factor of 0.4, and the parameter λ for Ridge regression was $\lambda = 1 \times 10^{-4}$. For each optimizer, we also recorded the mean squared error (MSE) at each iteration to study convergence.

The complete implementation is provided in Code Listing 6, and the resulting fits and convergence curves are shown in Figure 6.

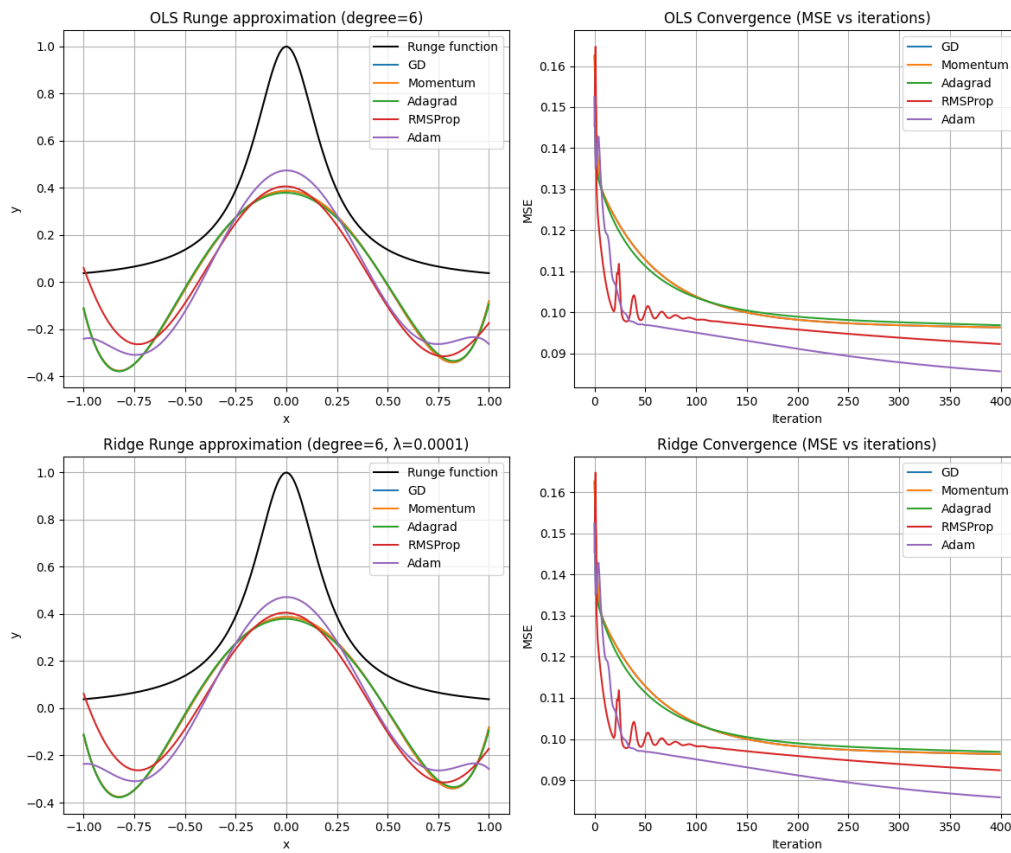


Figure 6: Optimizers fitted on the Runge Function

Part E: Writing own code for Lasso regression

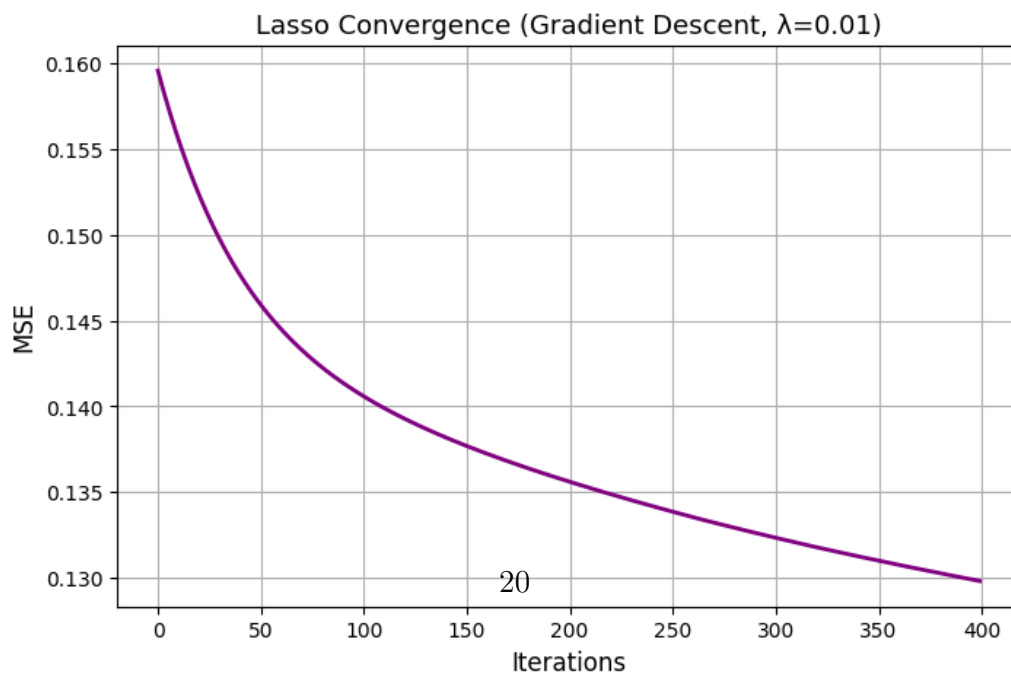


Figure 7: Lasso Convergence

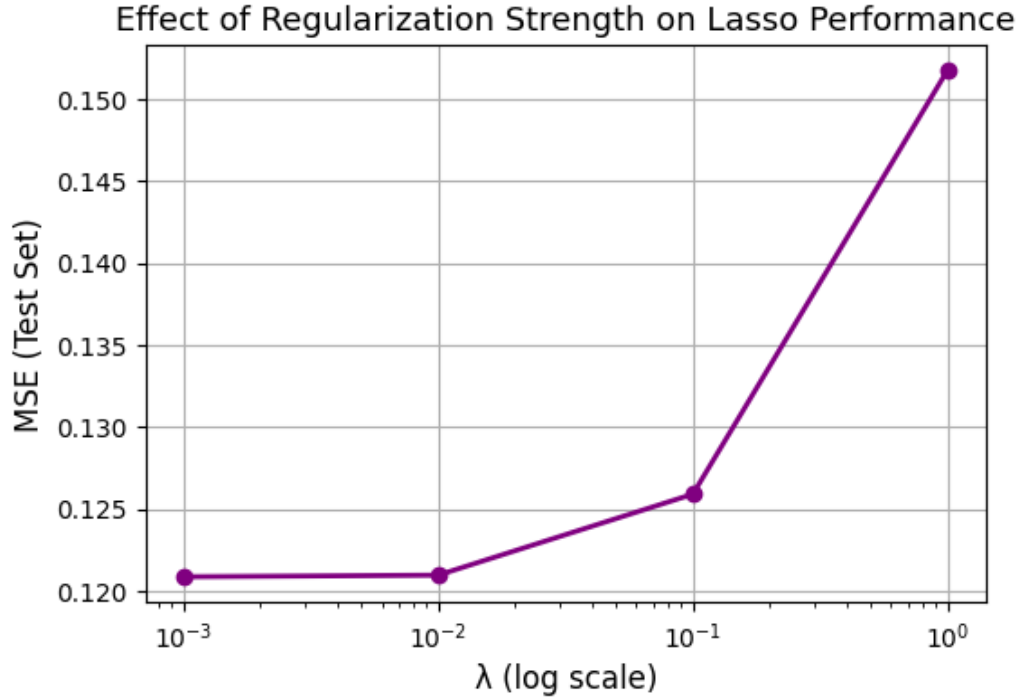


Figure 8: Effect of regularization strength on Lasso

Part F: Stochastic Gradient Descent

We implemented **SGD** and tested the adaptive optimizers **Momentum**, **AdaGrad**, **RMSProp**, and **Adam** for both **OLS** and **Ridge** regression on the Runge function. The models were trained using mini-batches with an 80/20 train-test split, a polynomial degree of 10, and a learning rate of $\eta = 0.03$. The results were compared with the baseline **vanilla gradient descent (GD)**.

For each optimizer, we recorded the mean squared error (MSE) at each iteration to study convergence behaviour. The implementation of the SGD procedure is provided in Code Listing 8, and the resulting convergence curves are shown in Figure 9.

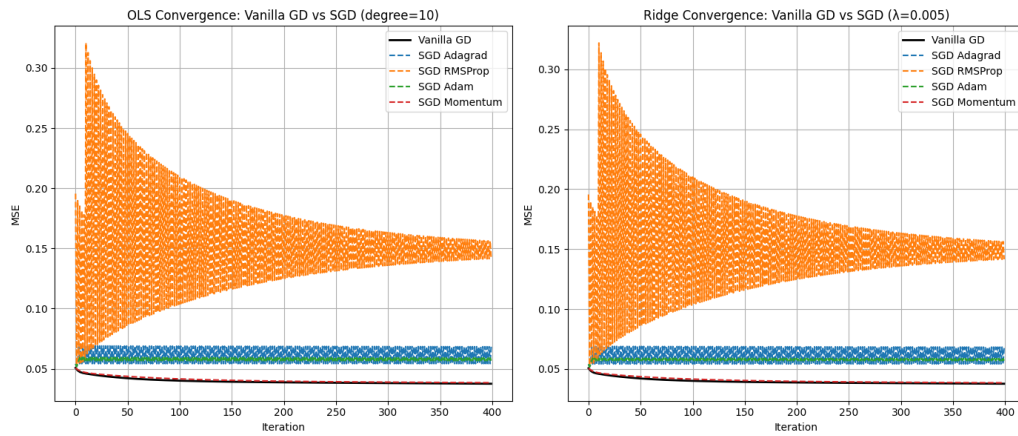


Figure 9: SGD convergence for OLS (left) and Ridge (right) regressions compared with vanilla gradient descent.

Part G: Bias-Variance Trade-off with Bootstrap Resampling




Figure 10: Total Error, Bias, and Variance, with 100 bootstrap resamplings for an Ordinary Least Squares (OLS) model on 100 data points.

Part H: Cross-Validation




Figure 11: 5-fold cross validation on our three linear regression models, optimized using fixed learning rate, full-batch gradient descent for 50 data points. Values of regularization hyperparameters specified in legend.

Discussion

Part A

As we can see from the results in figure 1, we observe that generally the **MSE** values decrease as and the R^2 values increase in the unscaled version with increasing polynomial degrees. This should not perplex us given that a more complicated model can account for more of the variation in the dataset than a more simple one. We can see that the lowest values for the error are given with a degree 6 and onward in the unscaled version of the **OLS**.

When observing the scaled data (2 bottom images): we can observe that at first glance we see more of a variation, but if we notice the scale of the data we can actually see that the variation of the data is quite significantly lower than in the unscaled versions. The scaling of the data magnifies the differences between the **MSE** values, making our plot less smooth, however we can see that having the scaled data makes makes our variance much lower. This demonstrates the effect of scaling on numerical stability and ties

directly to the bias–variance trade-off: scaling reduces variance and stabilizes coefficient estimates without changing the underlying model complexity.

Scaling ensures that all polynomial features contribute proportionally to the optimization process. Without scaling, higher-order terms dominate the regression due to their larger magnitudes, which leads to ill-conditioned matrices and unstable coefficient estimates. Standardizing the features (using `StandardScaler`) centers the data which improves numerical stability and convergence of the optimization. This stabilization reduces variance in the parameters but does not change the underlying model complexity as said above.

Part B

As we can see from the results in figure 2 we can see the scaled data for the ridge regression fitted for several different λ . As discussed in the section for **Ridge** regression, a lower value of λ will generally give us a lower **MSE** values due to penalization of the large coefficients being less. From the image we can see that the lowest error comes from the value $\lambda = 0.0001$. We can also see that in likeness to Part A, a more complex model will generally map out the parameters θ more optimally than a simpler model, however overfitting can happen.

To analyze further the effect of regularization, we computed the difference in MSE between Ridge and the corresponding OLS models across the dataset. The results, shown in Figure 3, confirm that larger λ values increase the overall error due to stronger coefficient shrinkage, while smaller λ values preserve model flexibility and yield lower MSE. This behavior reflects the fundamental bias–variance trade-off: as λ increases, variance decreases and bias increases. Hence, Ridge regression provides a mechanism to control model complexity and improve generalization compared to unregularized OLS. However as we see the scale of the λ values above 0.01 do not give any significant effect.

Part C

The graph in Figure 4 illustrates how Gradient Descent (GD) for OLS and Ridge regression behave differently in terms of Mean Squared Error (MSE) minimization. The optimization was performed with a fixed learning rate $\eta = 0.01$ over 10,000 iterations (as shown on the x -axis).

Both curves display a rapid initial decrease in MSE, indicating that the gradient descent algorithm quickly approaches a region close to the optimal solution.

The blue curve represents GD for OLS. It continues to decrease slowly toward very small MSE values, reflecting the fact that OLS lacks a regularization term and thus aims to minimize the training error as much as possible. However, this tendency to fully minimize the error may cause the model to overfit the training data, resulting in high variance.

In contrast, the orange curve represents GD for Ridge regression. Here, the MSE converges more quickly and stabilizes at a slightly higher value. This behavior arises from the inclusion of the regularization parameter λ , which penalizes large coefficients and effectively shrinks them toward zero. As a result, Ridge regression accepts a slightly higher bias (poorer fit on the training data) in exchange for lower variance and improved generalization performance on unseen data.

The graph in Figure 4 illustrates how Gradient Descent (GD) for OLS and Ridge regression behave differently in terms of Mean Squared Error (MSE) minimization. The optimization was performed with a fixed learning rate $\eta = 0.01$ over 10,000 iterations (as shown on the x -axis).

Both curves display a rapid initial decrease in MSE, indicating that the gradient descent algorithm quickly approaches a region close to the optimal solution.

The blue curve represents GD for OLS. It continues to decrease slowly toward very small MSE values, reflecting the fact that OLS lacks a regularization term and thus aims to minimize the training error as much as possible. However, this tendency to fully minimize the error may cause the model to overfit the training data, resulting in high variance.

In contrast, the orange curve represents GD for Ridge regression. Here, the MSE converges more quickly and stabilizes at a slightly higher value. This behavior arises from the inclusion of the regularization parameter λ , which penalizes large coefficients and effectively shrinks them toward zero. As a result, Ridge regression accepts a slightly higher bias (poorer fit on the training data) in exchange for lower variance and improved generalization performance on unseen data.

Part D

The implementation of our adaptive learning rate and momentum methods can be found in code listing 6. From the results in figure 6 we can observe

clear differences in the different methods. All optimizers were trained using a fixed learning rate of $\eta = 0.03$ and momentum factor 0.4, with $\lambda = 1 \times 10^{-4}$ for Ridge regression. The chosen polynomial degree was 6 since this proved to be the best one tested up to a degree of 15, where the program became slower and much more computationally expensive. Based on the results, the adaptive methods that mapped the Runge function best in order of performance were:

- 1) ADAM
- 2) RMSProp
- 3) Momentum
- 4) AdaGrad
- 5) Vanilla GD

The superiority of **Adam** is consistent with the theoretical expectations expressed in the section about gradient descent variants. By combining momentum and adaptive learning rate adjustments, the algorithm provides fast and stable convergence, as well as the best approximation.

RMSProp also performs well due to its exponentially weighted moving average of squared gradients, which helps maintain a stable learning rate. The algorithm is heavily influenced by the choice of hyperparameters chosen, and therefore might be better if it was tuned separately.

Momentum is an improvement above the vanilla gradient descent, as well as the AdaGrad, where it smoothes the path of the gradient descent and accelerates convergence if the direction is consistent.

AdaGrad, while stable, performs worse due to its continuously decaying learning rate, which eventually becomes too small for further updates. The lack of performance is probably due to the iteration number chosen in this task (400 iterations).

The overall results highlight how the adaptive learning rate algorithms improve the convergence and numerical stability. We have chosen to overlay the actual function with those created by utilizing the algorithms so we can visually get a sense of how well they have performed. Adam and RMSProp provide the best performance, while AdaGrad provides the worst performance because of its decaying step size limit. These findings are in line with the theoretical behaviors described in the [9].

0.1 Part E

Figure 7 shows the convergence of the LASSO regression model using gradient descent for $\lambda = 0.01$. The mean squared error (MSE) decreases steadily with

each iteration, demonstrating that the optimization algorithm is functioning correctly and approaching a minimum. The convergence is smooth and monotonic, indicating that the chosen learning rate and number of iterations are appropriate for stable optimization. The final MSE stabilizes around 0.13, suggesting that the model reaches a satisfactory trade-off between fitting the data and applying regularization.

Figure 8 illustrates how the test MSE varies with the regularization parameter λ . For small λ values (10^{-3} – 10^{-2}), the MSE remains low, implying that mild regularization helps control variance without significantly increasing bias. As λ increases beyond 0.1, the MSE rises sharply, indicating that excessive regularization causes underfitting by shrinking relevant coefficients too strongly towards zero. This behaviour aligns with the expected bias–variance trade-off inherent in LASSO regression.

Overall, the results confirm that LASSO effectively regularizes the model by penalizing large coefficients and promoting sparsity. However, for the smooth Runge function, strong LASSO regularization can remove useful higher-order polynomial terms, resulting in poorer generalization performance. The optimal performance is achieved for small λ values, where the model balances bias and variance effectively.

Compared to the gradient descent results for OLS and Ridge regression, the LASSO method exhibits slower convergence and higher final MSE values. This behavior arises from the non-smooth L_1 penalty, which makes the optimization problem more difficult and less stable near zero coefficients. While Ridge regression benefits from the smooth L_2 penalty that improves conditioning and stabilizes convergence, LASSO must rely on subgradient updates, resulting in a more gradual descent. Furthermore, the regularization path of LASSO shows that increasing λ causes several coefficients to shrink exactly to zero, producing a sparse model. Although this sparsity can be beneficial for feature selection, it introduces additional bias and leads to underfitting for the smooth Runge function, where many polynomial terms contribute meaningfully. In contrast, Ridge achieves a better bias–variance balance, whereas LASSO prioritizes simplicity at the expense of accuracy.

Part F

In this final part, we extended our gradient descent implementations by introducing **stochastic gradient descent (SGD)** using the same adaptive learning rate methods as in Parts C–E. The implementation is shown in Code Listing 8, and the results are presented in Figure 9. As is the practice

throughout this project: the train-test split was 80/20. We compare the convergence behaviour of SGD with that of the full-batch gradient descent used previously. In our SGD we introduce the mini batches and sampling of the dataset, which leads to higher variance in our dataset. Based on the results, the adaptive methods that had the least **MSE** values in order were:

- 1)Vanilla GD
- 2)Momentum SGD
- 3)ADAM SGD
- 4)AdaGrad SGD
- 5)RMSPProp SGD

Most interestingly, the vanilla GD achieved the lowest final MSE. This is perhaps because the full-batch gradient computes the exact gradient at each step, giving us a more stable and smoother conversion. We don't actually see any of the SGD implementations outperforming the vanilla GD with the varying adaptive learning rates. Perhaps this is a result of the parameters chosen, but also perhaps this is a flaw of our code.

Evidently from the results in Figure 9, the **Momentum** optimizer performed closest to the baseline **Vanilla GD**, achieving the best performance among the adaptive methods. **Adam** followed closely in third place, while **AdaGrad** ranked fourth with a noticeably higher variance in its MSE curve. This behaviour is consistent with theoretical expectations, as AdaGrad's learning rate decays over time, eventually becoming too small for significant updates.

Most surprisingly, the **RMSPProp** implementation performed the worst, displaying severe oscillations that far exceeded those of the other methods. This likely stems from the sensitivity of RMSPProp to hyperparameter tuning—particularly the decay factor ρ and the learning rate η . Even after testing multiple parameter configurations, the instability persisted, suggesting that the combination of stochastic mini-batches and our chosen parameters amplified oscillatory behaviour.

Overall, in our case, **SGD did not outperform vanilla GD** in terms of final convergence. While stochastic updates can improve per-iteration efficiency, they also introduce gradient noise that can hinder convergence to the true minimum, especially for smaller datasets or when the learning rate is not finely tuned. Momentum and Adam helped mitigate this noise, but

the fully deterministic gradient descent ultimately produced the lowest final MSE and smoothest convergence.

Part G

Figure 10 illustrates the bias–variance trade-off for an Ordinary Least Squares (OLS) model fitted to the noisy Runge function using 100 bootstrap resamplings. The plot shows the total error, bias, and variance as functions of polynomial degree. As the polynomial degree increases from 1 to 15, we observe the following trends:

- **Bias:** The bias decreases consistently with increasing polynomial degree. This is expected, as more complex models can better capture the underlying function, reducing systematic error.
- **Variance:** The variance increases with polynomial degree. Simpler models (low degree) are stable and have low variance, while complex models (high degree) are sensitive to fluctuations in the training data, leading to high variance.
- **Total Error:** The total error, which is the sum of bias and variance, exhibits a U-shaped curve. It decreases initially as bias reduction dominates, reaching a minimum between 6 and 12, before increasing again as variance becomes the dominant factor.

These results illustrate the classic bias–variance trade-off: simple models (low degree) have high bias but low variance, while complex models (high degree) have low bias but high variance. The optimal polynomial degree balances these two sources of error, minimizing the total prediction error. In this case, a polynomial degree of around 6 achieves the best trade-off, yielding the lowest total error on average across the bootstrap samples.

Part H

Figure 11 presents the results of 5-fold cross-validation for three linear regression models: Ordinary Least Squares (OLS), Ridge, and Lasso. Each model was optimized using full-batch gradient descent with a fixed learning rate (selected based on the largest eigenvalue of the Hessian matrix of the cost function) over 50 data points sampled from the noisy Runge function. The regularization hyperparameters for Ridge ($\lambda = 0.0001$) and Lasso ($\lambda = 0.001$) were selected based on prior experiments and because of the small dataset. The plot shows the mean squared error (MSE) on the test folds as a function of polynomial degree from 1 to 15. The results illustrate several key points:

- **OLS:** The OLS model exhibits a U-shaped curve, with low MSE at low polynomial degrees (between 4 and 13) due to high bias. As the degree increases, the MSE rises sharply due to overfitting and high variance.
- **Ridge:** The Ridge regression model shows a slightly more stable MSE curve. The regularization helps control variance, allowing the model to maintain lower MSE at higher polynomial degrees compared to OLS.
- **Lasso:** The Lasso model also demonstrates improved stability over OLS, with the most consistently low MSE, even as the polynomial degree increases.

These effects are more pronounced with fewer data points, where the risk of overfitting is higher.

Overall, the cross-validation results confirm that regularization techniques like Ridge and Lasso effectively mitigate overfitting compared to unregularized OLS. In this case, LASSO, in particular seems the most effective for this setup.

If we compare these results to those from the bootstrap analysis in Part G, we see consistent trends. Both methods highlight the bias–variance trade-off, with OLS suffering from high variance at complex polynomial degrees. Ridge and Lasso regularization reduce variance, leading to more stable performance across folds and bootstrap samples.

Conclusion

This project provided valuable insight into the fundamental principles of machine learning and regression analysis. By using the Runge function as a test case, we demonstrated how increasing polynomial complexity can lead to overfitting when data points are evenly spaced. Although high-degree polynomials appear capable of capturing complex patterns, they tend to generalize poorly due to excessive sensitivity to training data. Through this, we gained a deeper understanding of the bias–variance tradeoff, a central concept in machine learning that highlights the balance between a model’s accuracy and its stability. Regularization techniques such as Ridge and LASSO proved effective in mitigating overfitting by penalizing large coefficients and controlling model complexity. Comparing different optimization methods, we observed that Ridge regression achieved the best generalization performance, while Adam provided the fastest and most stable convergence among the gradient-based algorithms. These findings emphasize the importance of choosing both an appropriate regularization scheme and optimization strategy for achieving reliable predictive models. However, our analysis was limited to

one-dimensional data and polynomial features. Future work could extend this framework to multivariate datasets or alternative basis functions, such as splines or kernel expansions, to better assess generalization in higher dimensions. Overall, this project deepened our practical and theoretical understanding of how regression methods, regularization, and optimization interact to control overfitting and improve predictive performance—concepts that are foundational across modern machine learning.

Source Code

All source code is available from our GitHub repository: https://github.com/komandoyoko/Yoan_FYS_STK4155/tree/main/Project1

Sources

References

- [1] Brian Beers, *Regression: Definition, Analysis, Calculation, and Example*, Investopedia, 2025. Available: <https://www.investopedia.com/terms/r/regression.asp> [Accessed: 14-Sep-2025].
- [2] GeeksforGeeks, “Mean Squared Error (MSE),” <https://www.geeksforgeeks.org/maths/mean-squared-error/>, accessed October 2025.
- [3] XLSTAT, *Ordinary Least Squares Regression (OLS)*, Addinsoft, 2025. Available: <https://www.xlstat.com/solutions/features/ordinary-least-squares-regression-ols>. [Accessed: 14-Sep-2025].
- [4] GeeksForGeeks, *What is Ridge Regression?*, GeeksForGeeks, 2025. Available: <https://www.geeksforgeeks.org/machine-learning/what-is-ridge-regression/> [Accessed: 14-Sep-2025].
- [5] GeeksforGeeks, *What is Lasso Regression?*, available at: <https://www.geeksforgeeks.org/machine-learning/what-is-lasso-regression/>, accessed October 2025.
- [6] IBM Corporation, “Gradient Descent,” IBM THINK, <https://www.ibm.com/think/topics/gradient-descent>, accessed September 2025.

- [7] CompPhysics, *MachineLearning: Lecture Notes, Week 36*, GitHub repository, <https://github.com/CompPhysics/MachineLearning/blob/master/doc/LectureNotes/week36.ipynb>, accessed September 2025.
- [8] M. Hjorth-Jensen, “Optimization, the central part of any Machine Learning algorithm” University of Oslo, https://compphysics.github.io/MachineLearning/doc/LectureNotes/_build/html/chapteroptimization.html, accessed September 2025.
- [9] M. Hjorth-Jensen, “Stochastic Gradient Descent vs. Full Batch Gradient Descent,” in *Machine Learning Lecture Notes, Week 37*, University of Oslo, https://compphysics.github.io/MachineLearning/doc/LectureNotes/_build/html/week37.html#sgd-vs-full-batch-gd-convergence-speed-and-memory-comparison, accessed September 2025.
- [10] GeeksforGeeks, “R-squared (Coefficient of Determination),” <https://www.geeksforgeeks.org/maths/r-squared/>, accessed October 2025.
- [11] Scikit-learn Developers, “StandardScaler — Scikit-learn Documentation,” https://scikit-learn.org/stable/auto_examples/preprocessing/plot_all_scaling.html#plot-all-scaling-standard-scaler-section, accessed October 2025.

A Code Listings

```
1 #choosing datapoints between 1 and -1
2 n = 100
3 x = np.linspace (-1,1, n)
4
5
6 #From the Runge function: for every x, compute the
   corresponding y
7 y = 1/(1 + 25 * x**2)
8 #Adding noise (Gaussian) with mean 0 and variance 1: N(0,1)
9 noise = np.random.normal(0,1,n)
10 y_noisy = y + 0.1 * noise
11
12
13 #Creating feature matrix with shape [1,x,x^2,x^3,x^4 ...]
14 def polynomial_features(x, p, intercept = True):
15     n = len(x)
```



```

16     if intercept == True: # adding functionality for
intercept
17         X = np.zeros((n, p + 1))
18         X[:, 0] = 1 # intercept column first
19         for i in range(1, p + 1): # from 1 to p inclusive
20             X[:, i] = x ** i
21     else: #if intercept is False
22         X = np.zeros((n, p)) #no intercept column
23         for i in range(p): #from 0 to p-1
24             X[:, i] = x ** (i + 1)
25
26     return X
27
28 def OLS_parameters(X, y):
29     X_T_X = np.linalg.pinv(X.T @ X)
30     beta_OLS = X_T_X @ X.T @ y
31     return beta_OLS
32
33 def R2_score(actual , predicted):
34     numerator = np.sum((actual - predicted) ** 2) #the
numerator of the r2 formula
35     actual_mean = np.mean(actual) #our mean value of the
actual values
36     denominator = np.sum((actual - actual_mean) ** 2) #
denominator of the r2 formula
37     R2 = 1 - (numerator / denominator) #calculating the r2
38     return R2
39
40
41 import matplotlib.pyplot as plt
42 from sklearn.preprocessing import StandardScaler
43
44 mse_values_OLS_unscaled , R2_values_OLS_unscaled = np.zeros
(15) ,np.zeros(15) #storing the mse values for each
polynomial degree - unscaled
45 mse_values_OLS_scaled , R2_values_OLS_scaled = np.zeros(15) ,
np.zeros(15) #storing the r2 values for each polynomial -
Scaled
46
47 for i in range(1 , 16):
48     #unscaled version
49     X = polynomial_features(x , i)
50     X_train , X_test , y_train , y_test = train_test_split(X
, y_noisy , test_size = 0.2 , random_state=67) #splitting
into train_test_split
51     beta = OLS_parameters(X_train , y_train) #finding optimal
beta
52     y_pred = X_test @ beta #predicting using the training
model

```

```

53     mse_values_OLS_unscaled[i-1] = mean_squared_error(y_test
54     , y_pred)
55     R2_values_OLS_unscaled[i-1] = R2_score(y_test , y_pred)
56
57     #scaled version
58     '''
59     Important note: We split, and then introduce the scaler
60     because we do not want to contaminate the dataset of train
61     with the information from test!!!
62     '''
63     scaler = StandardScaler()
64     X_train_scaled = scaler.fit_transform(X_train)
65     X_test_scaled = scaler.transform(X_test)
66
67     beta_scaled = OLS_parameters(X_train_scaled , y_train)
68     y_pred_scaled = X_test_scaled @ beta_scaled
69
70     mse_values_OLS_scaled[i-1] = mean_squared_error(y_test ,
71     y_pred_scaled)
72     R2_values_OLS_scaled[i-1] = R2_score(y_test ,
73     y_pred_scaled)
74
75     #now we create degrees array as well
76     degrees = np.arange(1 , 16)
77     plt.figure(figsize = (12 , 10))
78
79     # Unscaled MSE
80     plt.subplot(2 , 2 , 1)
81     plt.plot(degrees , mse_values_OLS_unscaled , marker = 'o',
82     color='red')
83     plt.xlabel('Polynomial Degree')
84     plt.ylabel('Mean Squared Error')
85     plt.title('MSE vs Polynomial Degree (Unscaled)')
86     plt.grid()
87
88     # Unscaled  $R^2$ 
89     plt.subplot(2 , 2 , 2)
90     plt.plot(degrees , R2_values_OLS_unscaled , marker = 'o',
91     color='orange')
92     plt.xlabel('Polynomial Degree')
93     plt.ylabel('R2 Score')
94     plt.title('R2 vs Polynomial Degree (Unscaled)')
95     plt.grid()
96
97     # Scaled MSE

```

```

95 plt.subplot(2 , 2 , 3)
96 plt.plot(degrees , mse_values_OLS_scaled , marker = 's',
    color='blue')
97 plt.xlabel('Polynomial Degree')
98 plt.ylabel('Mean Squared Error')
99 plt.title('MSE vs Polynomial Degree (Scaled)')
100 plt.grid()
101
102 # Scaled R^2
103 plt.subplot(2 , 2 , 4)
104 plt.plot(degrees , R2_values_OLS_scaled , marker = 's', color
    ='green')
105 plt.xlabel('Polynomial Degree')
106 plt.ylabel('R^2 Score')
107 plt.title('R^2 vs Polynomial Degree (Scaled)')
108 plt.grid()
109
110 plt.tight_layout()
111 plt.show()

```

Listing 1: Implementation of OLS regression

```

1 def R2_score(actual , predicted):
2     numerator = np.sum((actual - predicted) ** 2) #the
    numerator of the r2 formula
3     actual_mean = np.mean(actual) #our mean value of the
    actual values
4     denominator = np.sum((actual - actual_mean) ** 2) #
    denominator of the r2 formula
5     R2 = 1 - (numerator / denominator) #calculating the r2
6     return R2

```

Listing 2: R^2 Implementation

```

1 def ridge_regression(X , y , lamb):
2 def ridge_regression(X , y , lamb):
3     I = np.identity(X.shape[1]) # need to have the identity
    matrix
4     return (np.linalg.inv(X.T @ X + lamb * I) @ X.T @ y) #the
    actual expression for optimal beta
5
6 #setting differing values for lambda
7 lambd = [0.0001 , 0.01 , 0.1 , 0.5]
8
9 mse_values_ridge = np.zeros((15 , len(lambd))) #defining the
    mse values matrix
10 R2_values_ridge = np.zeros((15, len(lambd))) #R2 matrix for
    storage
11
12 for l in range(len(lambd)): #we want to test for all lambdas

```

```

13     for i in range(1 , 16): #we want to test for teh
different polynomials
14         X = polynomial_features(x , i)
15         X_train , X_test , y_train , y_test =
train_test_split(X , y_noisy , test_size = 0.2 ,
random_state= 60) #split the data
16
17         scaler = StandardScaler()
18         X_train = scaler.fit_transform(X_train)
19         X_test = scaler.transform(X_test)
20
21         beta_ridge = ridge_regression(X_train , y_train ,
lambda[1]) #do the ridge regression
22         y_pred_ridge = X_test @ beta_ridge #do the prediction
23
24         mse_values_ridge[i-1, 1 ] = mean_squared_error(y_test
, y_pred_ridge) #calculate the mse values
25         R2_values_ridge[i-1 , 1] = R2_score(y_test ,
y_pred_ridge) #calculate the r2 values
26
27
28
29
30 plt.figure(figsize=(12, 5))
31
32 # ---- MSE plot ----
33 plt.subplot(1, 2, 1)
34 for l in range(len(lambda)):
35     plt.plot(degrees, mse_values_ridge[:, 1], marker='o',
label=f"\lambda={lambda[l]}")
36 plt.xlabel('Polynomial Degree')
37 plt.ylabel('Mean Squared Error')
38 plt.title('MSE Ridge vs Polynomial Degree')
39 plt.legend()
40 plt.grid()
41
42 # ---- R^2 plot ----
43 plt.subplot(1, 2, 2)
44 for l in range(len(lambda)):
45     plt.plot(degrees, R2_values_ridge[:, 1], marker='o',
label=f"\lambda={lambda[l]}")
46 plt.xlabel('Polynomial Degree')
47 plt.ylabel('R^2 Score')
48 plt.title('R^2 Ridge vs Polynomial Degree')
49 plt.legend()
50 plt.grid()
51
52 plt.tight_layout()

```

```
53 plt.show()
```

Listing 3: Ridge Implementation

```
1 difference_mse = np.zeros((15 , 4))
2 difference_r2 = np.zeros((15 , 4))
3 for j in range(len(lambd)):
4     for i in range(1 , 16):
5         difference_mse[i-1, j] = mse_values_ridge[i-1, j] -
mse_values_OLS_scaled[i-1]
6         difference_r2[i-1 , j] = R2_values_ridge[i-1 , j] -
R2_values_OLS_scaled[i-1]
7
8
9 mean_diff_mse = np.mean(difference_mse, axis=0)    # one value
per \lambda
10 mean_diff_r2  = np.mean(difference_r2, axis=0)    # one value
per \lambda
11
12 plt.figure(figsize=(10,4))
13
14 # ---- Mean difference in MSE ----
15 plt.subplot(1,2,1)
16 plt.bar([str(l) for l in lambd], mean_diff_mse, color='
skyblue')
17 plt.axhline(0, color='black', linestyle='--')
18 plt.xlabel("\lambda")
19 plt.ylabel("Mean \Delta MSE (Ridge - OLS)")
20 plt.title("Average MSE Difference")
21
22 # ---- Mean difference in R^2 ----
23 plt.subplot(1,2,2)
24 plt.bar([str(l) for l in lambd], mean_diff_r2, color='salmon'
)
25 plt.axhline(0, color='black', linestyle='--')
26 plt.xlabel("\lambda")
27 plt.ylabel("Mean R^2 (Ridge - OLS)")
28 plt.title("Average R^2 Difference")
29
30 plt.tight_layout()
31 plt.show()
```

Listing 4: Capturing difference between OLS and Ridge using MSE and R^2

```
1 eta = 0.01
2 num_iters = 10000
3 lam = 0.1
4
```

```

5 n_samples, n_features = X.shape
6
7 # Initialize parameters to zero
8 theta_gdOLS = np.zeros(n_features)
9 theta_gdRidge = np.zeros(n_features)
10
11 def ols_gradient(X, y, theta):
12     return (2/n_samples) * X.T @ (X @ theta - y)
13
14 def ridge_gradient(X, y, theta, lam):
15     return (2/n_samples) * X.T @ (X @ theta - y) + 2 * lam *
    theta
16
17 # Track MSE
18 ols_loss = []
19 ridge_loss = []
20
21 for t in range(num_iters):
22     grad_OLS = ols_gradient(X, y, theta_gdOLS)
23     grad_Ridge = ridge_gradient(X, y, theta_gdRidge, lam)
24
25     theta_gdOLS -= eta * grad_OLS
26     theta_gdRidge -= eta * grad_Ridge
27
28     ols_loss.append(np.mean((y - X @ theta_gdOLS) ** 2))
29     ridge_loss.append(np.mean((y - X @ theta_gdRidge) ** 2))
30
31 print("Gradient Descent OLS coefficients:", theta_gdOLS)
32 print("Gradient Descent Ridge coefficients:", theta_gdRidge)
33
34 #Closed-form solution
35 def OLS_parameters(X, y):
36     return np.linalg.inv(X.T @ X) @ X.T @ y
37
38 def Ridge_parameters(X, y, lam):
39     I = np.eye(X.shape[1])
40     return np.linalg.inv(X.T @ X + lam * I) @ X.T @ y
41
42 theta_closed_formOLS = OLS_parameters(X, y)
43 theta_closed_formRidge = Ridge_parameters(X, y, lam)
44
45 print("Closed-form OLS:", theta_closed_formOLS)
46 print("Closed-form Ridge:", theta_closed_formRidge)
47
48 #plot GD Convergence
49 plt.figure(figsize=(10, 4))
50 plt.plot(range(num_iters), ols_loss, label='OLS (GD)')
51 plt.plot(range(num_iters), ridge_loss, label='Ridge (GD)')
52 plt.xlabel('Iterations')

```

```

53 plt.ylabel('MSE')
54 plt.title('Gradient Descent Convergence')
55 plt.legend()
56 plt.grid(True, linestyle='--', alpha=0.7)
57 plt.tight_layout()
58 plt.show()
59
60
61 # Plot Model Predictions vs True Function
62 y_pred_OLS_gd      = X @ theta_gdOLS
63 y_pred_Ridge_gd     = X @ theta_gdRidge
64 y_pred_OLS_closed   = X @ theta_closed_formOLS
65 y_pred_Ridge_closed = X @ theta_closed_formRidge
66
67 plt.figure(figsize=(10, 6))
68 plt.scatter(x, y, color='steelblue', alpha=0.6, label='Noisy
    data')
69 plt.plot(x, y, color='black', linewidth=2, label='True Runge
    function')
70 plt.plot(x, y_pred_OLS_closed, '--', color='darkblue', label=
    'Closed-form OLS')
71 plt.plot(x, y_pred_OLS_gd, ':', color='orange', label='GD OLS
    ')
72 plt.plot(x, y_pred_Ridge_closed, '--', color='green', label='
    Closed-form Ridge')
73 plt.plot(x, y_pred_Ridge_gd, ':', color='red', label='GD
    Ridge')
74 plt.xlabel('x')
75 plt.ylabel('y')
76 plt.title('Model Predictions vs True Function')
77 plt.legend()
78 plt.grid(True, linestyle='--', alpha=0.7)
79 plt.tight_layout()
80 plt.show()

```

Listing 5: Gradient descent for OLS and Ridge

```

1 import numpy as np
2 from sklearn.metrics import mean_squared_error
3 from sklearn.model_selection import train_test_split
4 from sklearn.preprocessing import PolynomialFeatures,
    StandardScaler
5 import matplotlib.pyplot as plt
6
7
8 n_samples = 200 #number of samples we will be using
9
10
11 def ols_gradient(X, y, theta):

```

```

12     return (2/n_samples) * X.T @ (X @ theta - y)
13
14 def ridge_gradient(X, y, theta, lam):
15     return (2/n_samples) * X.T @ (X @ theta - y) + 2 * lam *
    theta
16
17
18
19
20
21 def momentum_gd(X , y , iterations, momentum , n_steps , func
    , lam , theta_init = None):
22     if theta_init is None:
23         theta = np.zeros(X.shape[1])
24     else:
25         theta = theta_init.copy() #we start by implementing
    the starting point of the theta
26
27     change = np.zeros_like(theta) #we initialize the change
28
29     mse_val = np.zeros(iterations) # We also initialize
    mse_val as the measure of how far away from our goal we
    are
30     for i in range(iterations): #here we want to iterate
    until we reach our point
31         if func == "OLS": #we split between OLS and ridge
    gradients
32             grad = ols_gradient(X , y , theta)
33         elif func == "Ridge":
34             grad = ridge_gradient(X , y , theta , lam)
35
36         change = momentum * change + n_steps * grad #the
    change is based on the momentum * change + direction we
    are headed
37         theta -= change #the final result is theta is changed
    based on the previous change + the new one.
38
39         mse_val[i] = mean_squared_error(y , X @ theta) #we
    want to calculate the error away from the true values of
    the function
40     return theta , mse_val
41
42
43
44
45
46 def ADAGRAD(X , y , iterations , n_steps , func , lam , eps
    = 1e-6 , theta_init = None):
47     if theta_init is None: #if we do not have an initial

```



```

48     theta = np.zeros(X.shape[1])
49     else:
50         theta = theta_init.copy() #initialize the parameters
we need from the function call
51     r = np.zeros_like(theta)
52
53     mse_val = np.zeros(iterations)
54     for i in range(iterations):
55         if func == "OLS":
56             grad = ols_gradient(X , y , theta)
57         elif func == "Ridge":
58             grad = ridge_gradient(X , y , theta , lam)
59
60         gtgt = grad ** 2 #here we are taking the square of
the gradient
61         r= r+ gtgt #we apply the adaptive step
62
63         theta = theta - (n_steps / (np.sqrt(r) + eps)) * grad
#we calculate the new theta here based on the adaptive
step
64
65         mse_val[i] = mean_squared_error(y , X @ theta) #and
we also calculate the MSE
66     return theta , mse_val
67
68
69
70
71 def RMSProp(X , y , iterations , n_steps , func , lam , p =
0.9 , eps = 1e-5, theta_init = None):
72     if theta_init is None:
73         theta = np.zeros(X.shape[1])
74     else:
75         theta = theta_init.copy() #inititalize the first steps
76     v = np.zeros_like(theta)
77
78     mse_val = np.zeros(iterations)
79     for i in range(iterations):
80         if func == "OLS":
81             grad = ols_gradient(X , y , theta)
82         elif func == "Ridge":
83             grad = ridge_gradient(X , y , theta , lam)
84
85         v = p*v + (1-p) * grad**2 # Here is our step
86
87         theta = theta - (n_steps/np.sqrt(v + eps)) * grad #
the new theta calculated from the RMSprop algorythm
88

```

```

89         mse_val[i] = mean_squared_error(y , X @ theta) #as
always we calculate the MSE
90     return theta , mse_val
91
92
93
94 def ADAM(X , y , iterations , n_steps , func , lam , b1 =
0.9 , b2 = 0.999 , eps = 1e-5, theta_init = None):
95     if theta_init is None:
96         theta = np.zeros(X.shape[1])
97     else:
98         theta = theta_init.copy() # we initialize the
different things we need
99     v = np.zeros_like(theta)
100    m = np.zeros_like(theta)
101
102    mse_val = np.zeros(iterations)
103    for i in range(1 , iterations + 1): #we cannot start with
0 in this loop because of the bias corrections need to
be raised to the power of the iteration.
104        if func == "OLS":
105            grad = ols_gradient(X , y , theta)
106        elif func == "Ridge":
107            grad = ridge_gradient(X , y , theta , lam)
108
109        m = b1*m + (1-b1) * grad #we define our unbiased
terms
110        v = b2*v + (1-b2)* grad ** 2
111
112        m_correct = m / (1-b1**i) #here we apply the bias
correction terms
113        v_correct = v / (1-b2**i)
114
115        theta = theta - ( n_steps / (np.sqrt(v_correct) + eps
) ) * m_correct #calculate the theta according to the
algorythm
116
117        mse_val[i-1] = mean_squared_error(y , X @ theta) #
calculcate the MSE as always
118    return theta , mse_val
119
120
121
122
123
124
125
126
127

```

```

128
129
130
131 '''
132 In this block of the code we define the Runge function,
    create a design matrix and try to extract the optimal
    theta values using the
133 gradient descents we have implemented. This will be plotted
    against the actual runge function to visually see how
    close we are, and we will
134 also implement a plot of the MSE vs iterations of the
    different methods
135 '''
136 x = np.linspace(-1, 1, 200) #we define x
137 y_true = 1 / (1 + 25 * x**2) #the true runge function
138 lam = 1e-4
139 degree = 6 #we choose a polynomial of 4th degree
140
141
142 # setup
143 poly = PolynomialFeatures(degree)
144 X = poly.fit_transform(x.reshape(-1, 1))
145
146 scaler = StandardScaler()
147 X_scaled = scaler.fit_transform(X) # mean/std learned here
148
149 X_train, X_test, y_train, y_test = train_test_split(X_scaled,
    y_true, test_size=0.2) #We then split the data
150
151
152
153
154
155
156
157 '''
158 Here we define a function called run_optimizer that will take
    all the different functions and go through them creating
    different gradient descent methods
159 and extracting the mse_val and theta values
160 '''
161
162
163 def run_optimizer_OLS(method, X, y, lam):
164     iterations = 400
165     n_steps = 0.03
166     momentum = 0.4
167     lam = lam
168     if method == "GD":

```

```

169         # Standard gradient descent (no momentum)
170         return momentum_gd(X, y, iterations, momentum,
n_steps, "OLS" , lam)
171     elif method == "Momentum":
172         return momentum_gd(X, y, iterations, momentum,
n_steps, "OLS" , lam)
173     elif method == "Adagrad":
174         return ADAGrad(X, y, iterations, n_steps, "OLS" , lam
)
175     elif method == "RMSProp":
176         return RMSProp(X, y, iterations, n_steps, "OLS" , lam
)
177     elif method == "Adam":
178         return ADAM(X, y, iterations, n_steps, "OLS" , lam)
179     else:
180         raise ValueError("Unknown method")
181
182 def run_optimizer_ridge(method, X, y, lamb):
183     iterations = 400
184     n_steps = 0.03
185     momentum = 0.4
186     lam = lamb
187     if method == "GD":
188         return momentum_gd(X, y, iterations, momentum,
n_steps, "Ridge", lam)
189     elif method == "Momentum":
190         return momentum_gd(X, y, iterations, momentum,
n_steps, "Ridge", lam)
191     elif method == "Adagrad":
192         return ADAGrad(X, y, iterations, n_steps, "Ridge",
lam)
193     elif method == "RMSProp":
194         return RMSProp(X, y, iterations, n_steps, "Ridge",
lam)
195     elif method == "Adam":
196         return ADAM(X, y, iterations, n_steps, "Ridge", lam)
197     else:
198         raise ValueError("Unknown method")
199
200
201
202
203
204
205
206
207
208
209

```

```

210
211
212 if __name__ == "__main__" :
213
214     plt.figure(figsize=(12, 10))
215
216     # --- OLS fit ---
217     plt.subplot(2, 2, 1)
218     plt.plot(x, y_true, 'k', label="Runge function")
219     for method in ["GD", "Momentum", "Adagrad", "RMSProp", "
Adam"]:
220         theta, mse_val = run_optimizer_OLS(method, X_train,
y_train, lam)
221
222         # Use the same poly + scaler from above
223         X_pred = poly.transform(x.reshape(-1, 1))
224         X_pred = scaler.transform(X_pred) # <--- FIX: use
transform, not fit_transform
225         y_pred = X_pred @ theta
226
227         plt.plot(x, y_pred, label=method)
228     plt.legend()
229     plt.title(f"OLS Runge approximation (degree={degree})")
230     plt.xlabel("x")
231     plt.ylabel("y")
232     plt.grid(True)
233
234
235
236
237     # OLS convergence
238     plt.subplot(2, 2, 2)
239     for method in ["GD", "Momentum", "Adagrad", "RMSProp", "
Adam"]:
240         _, mse_val = run_optimizer_OLS(method, X_train,
y_train, lam)
241         plt.plot(mse_val, label=method)
242     plt.legend()
243     plt.title("OLS Convergence (MSE vs iterations)")
244     plt.xlabel("Iteration")
245     plt.ylabel("MSE")
246     plt.grid(True)
247
248
249
250
251
252
253

```

```

254
255
256
257
258
259     # Ridge fit
260     plt.subplot(2, 2, 3)
261     plt.plot(x, y_true, 'k', label="Runge function")
262     for method in ["GD", "Momentum", "Adagrad", "RMSProp", "
Adam"]:
263         theta, mse_val = run_optimizer_ridge(method, X_train,
y_train, lam)
264
265         X_pred = poly.transform(x.reshape(-1, 1))
266         X_pred = scaler.transform(X_pred)    # <--- same
scaler here too
267         y_pred = X_pred @ theta
268
269         plt.plot(x, y_pred, label=method)
270     plt.legend()
271     plt.title(f"Ridge Runge approximation (degree={degree}, \
lambda={lam})")
272     plt.xlabel("x")
273     plt.ylabel("y")
274     plt.grid(True)
275
276     # Ridge convergence
277     plt.subplot(2, 2, 4)
278     for method in ["GD", "Momentum", "Adagrad", "RMSProp", "
Adam"]:
279         _, mse_val = run_optimizer_ridge(method, X_train,
y_train, lam)
280         plt.plot(mse_val, label=method)
281     plt.legend()
282     plt.title("Ridge Convergence (MSE vs iterations)")
283     plt.xlabel("Iteration")
284     plt.ylabel("MSE")
285     plt.grid(True)
286
287     plt.tight_layout()
288     plt.show()

```

Listing 6: Different adaptive learning rates

```

1 import numpy as np
2 from sklearn.metrics import mean_squared_error
3 from project1_d import ADAGRAD, RMSProp, ADAM
4
5 def lasso_gradient(X, y, theta, lam):

```

```

6     n_samples = X.shape[0]
7     grad = (2 / n_samples) * X.T @ (X @ theta - y)
8     subgrad = lam * np.sign(theta)
9     subgrad[0] = 0.0
10    return grad + subgrad
11
12    def run_optimizer_lasso(method, X, y, lam=0.01):
13        iterations = 400
14        n_steps = 0.001
15        momentum = 0.9
16
17        if method in ["GD", "Momentum"]:
18            theta = np.zeros(X.shape[1])
19            change = np.zeros_like(theta)
20            mse_val = np.zeros(iterations)
21            mom = 0 if method == "GD" else momentum
22
23            for i in range(iterations):
24                grad = lasso_gradient(X, y, theta, lam)
25                change = mom * change - n_steps * grad
26                theta += change
27                mse_val[i] = mean_squared_error(y, X @ theta)
28
29            return theta, mse_val
30
31        elif method == "Adagrad":
32            return ADAGrad(X, y, iterations, n_steps, "Lasso",
33                           lam)
34        elif method == "RMSProp":
35            return RMSProp(X, y, iterations, n_steps, "Lasso",
36                           lam)
37        elif method == "Adam":
38            return ADAM(X, y, iterations, n_steps, "Lasso", lam)
39        else:
40            raise ValueError("Unknown method")
41
42    # Part E : plotting
43    #regularization strengths to test
44    lambdas = [0.001, 0.01, 0.1, 1]
45
46    # Store results for each
47    results = {}
48
49    print("=== Lasso Regression (Gradient Descent) ===")
50    for lam in lambdas:
51        theta, mse_val = run_optimizer_lasso("GD", X_train,
52                                              y_train, lam=lam)
53        y_pred = X_test @ theta

```

```

52     mse = mean_squared_error(y_test, y_pred)
53     r2 = r2_score(y_test, y_pred)
54     results[lam] = (mse_val, mse, r2)
55
56     print(f"    ={lam:<6}: MSE={mse:.4f}, R    ={r2:.4f}")
57
58 # Plot MSE convergence for one example
59 example_lambda = 0.01
60 plt.figure(figsize=(8, 5))
61 plt.plot(results[example_lambda][0], color='purple',
62         linewidth=2)
63 plt.xlabel('Iterations', fontsize=12)
64 plt.ylabel('MSE', fontsize=12)
65 plt.title(f'Lasso Convergence (Gradient Descent,    ={
66     example_lambda})', fontsize=13)
67 plt.grid(True)
68 plt.show()
69
70 # Plot MSE vs
71 mse_values = [results[lam][1] for lam in lambdas]
72
73 plt.figure(figsize=(6, 4))
74 plt.plot(lambdas, mse_values, marker='o', color='purple',
75         linewidth=2)
76 plt.xscale('log')
77 plt.xlabel('    (log scale)', fontsize=12)
78 plt.ylabel('MSE (Test Set)', fontsize=12)
79 plt.title('Effect of Regularization Strength on Lasso
80     Performance', fontsize=13)
81 plt.grid(True)
82 plt.show()

```

Listing 7: Lasso implementation and plots

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3 from sklearn.preprocessing import PolynomialFeatures ,
4     StandardScaler
5
6
7
8 from project1_d import ADAgrad , RMSProp , ADAM ,
9     momentum_gd
10
11
12
13 def Stochastic_gradient(X, y, lambd, method, optimizer,
14     size_batch, epochs, total_updates=400):
15     datapoint = len(y)
16     n_steps = 0.03 # learning rate
17     momentum = 0.3 # momentum factor
18     lam = lambd

```



```

13     theta = np.zeros(X.shape[1])
14
15     mse_vals = []
16     updates = 0    # keep track of total updates
17
18     while updates < total_updates:    # run until we have the
19         same number of updates as vanilla GD
20         indices = np.random.permutation(datapoint) #running a
21         random on the datapoints to shuffle
22         X_shuffled, y_shuffled = X[indices], y[indices] #
23         createing the new shuffled versions
24
25         for i in range(0, datapoint, size_batch):
26             if updates >= total_updates:
27                 break    # stop exactly at total_updates
28
29             X_batch = X_shuffled[i:i+size_batch] #creating
30             the minibatches
31             y_batch = y_shuffled[i:i+size_batch]
32
33             if method == "OLS" and optimizer == "Adagrad":
34                 theta, mse_val = ADAgrad(X_batch, y_batch,
35                 iterations=1, n_steps=n_steps, func="OLS", lam=lam,
36                 theta_init=theta)
37             elif method == "OLS" and optimizer == "RMSProp":
38                 theta, mse_val = RMSProp(X_batch, y_batch,
39                 iterations=1, n_steps=n_steps, func="OLS", lam=lam,
40                 theta_init=theta)
41             elif method == "OLS" and optimizer == "Adam":
42                 theta, mse_val = ADAM(X_batch, y_batch,
43                 iterations=1, n_steps=n_steps, func="OLS", lam=lam,
44                 theta_init=theta)
45             elif method == "OLS" and optimizer == "Momentum":
46                 theta, mse_val = momentum_gd(X_batch, y_batch
47                 , iterations=1, momentum=momentum, n_steps=n_steps, func="
48                 OLS", lam=lam, theta_init=theta)
49
50             elif method == "Ridge" and optimizer == "Adagrad"
51             :
52                 theta, mse_val = ADAgrad(X_batch, y_batch,
53                 iterations=1, n_steps=n_steps, func="Ridge", lam=lam,
54                 theta_init=theta)
55             elif method == "Ridge" and optimizer == "RMSProp"
56             :
57                 theta, mse_val = RMSProp(X_batch, y_batch,
58                 iterations=1, n_steps=n_steps, func="Ridge", lam=lam,
59                 theta_init=theta)
60             elif method == "Ridge" and optimizer == "Adam":
61                 theta, mse_val = ADAM(X_batch, y_batch,

```

```

iterations=1, n_steps=n_steps, func="Ridge", lam=lam,
theta_init=theta)
44         elif method == "Ridge" and optimizer == "Momentum"
":
45             theta, mse_val = momentum_gd(X_batch, y_batch
, iterations=1, momentum=momentum, n_steps=n_steps, func="
Ridge", lam=lam, theta_init=theta)
46
47             # store MSE
48             mse_vals.append(mse_val[-1] if isinstance(mse_val
, (list, np.ndarray)) else mse_val) #Here we append the
MSE values, and create a robust append method
49
50             updates += 1
51
52         return theta, np.array(mse_vals)
53
54
55
56
57 x = np.linspace(-3 , 3 , 80)
58 y = 1/(1+25*x**2)
59
60 lam = 0.005
61 degree = 10
62
63 X = PolynomialFeatures(degree).fit_transform(x.reshape(-1 ,
1)) #here we create the polynomial features
64 X = StandardScaler().fit_transform(X) #scale the data
65
66
67
68
69 from project1_d import run_optimizer_OLS ,
run_optimizer_ridge #import the optimizers ran in previous
assignment
70
71
72
73 optimizers = [ "Adagrad", "RMSProp", "Adam", "Momentum"] #
here we call our optimizers for the adaptive learning step
74
75 fig, axes = plt.subplots(1, 2, figsize=(14, 6))
76
77 # --- OLS ---
78 theta_gd, mse_gd = run_optimizer_OLS("GD", X, y, lambd=0)
79 axes[0].plot(mse_gd, label="Vanilla GD", color='black',
linewidth=2)
80

```

```

81 for opt in optimizers:
82     theta, mse_val = Stochastic_gradient(
83         X, y,
84         lamdb=0,
85         method="OLS",
86         optimizer=opt,
87         size_batch=80,
88         epochs=400
89     )
90     axes[0].plot(mse_val, label=f"SGD {opt}", linestyle='--')
91
92 axes[0].set_title(f"OLS Convergence: Vanilla GD vs SGD (
93     degree={degree})")
94 axes[0].set_xlabel("Iteration")
95 axes[0].set_ylabel("MSE")
96 axes[0].legend()
97 axes[0].grid(True)
98
99 # --- Ridge ---
100 theta_gd, mse_gd = run_optimizer_ridge("GD", X, y, lamb=lamb)
101 axes[1].plot(mse_gd, label="Vanilla GD", color='black',
102     linewidth=2)
103
104 for opt in optimizers:
105     theta, mse_val = Stochastic_gradient(
106         X, y,
107         lamdb=lamb,
108         method="Ridge",
109         optimizer=opt,
110         size_batch=80,
111         epochs=400
112     )
113     axes[1].plot(mse_val, label=f"SGD {opt}", linestyle='--')
114
115 axes[1].set_title(f"Ridge Convergence: Vanilla GD vs SGD (
116     lambda={lamb})")
117 axes[1].set_xlabel("Iteration")
118 axes[1].set_ylabel("MSE")
119 axes[1].legend()
120 axes[1].grid(True)
121
122 plt.tight_layout()
123 plt.show()

```

Listing 8: Stochastic Gradient Descent

B Object-oriented approach

To improve code organization and reusability, we experimented with encapsulating the regression methods within a class structure. This proved to be useful for managing the added complexity of the project, especially when dealing with multiple regression techniques and optimization algorithms. We used this approach when implementing bootstrap resampling and cross-validation in the last part of the project.

The proposed structure is based on a base OLS class, with Ridge and Lasso classes inheriting from it. They share common methods for fitting the model and making predictions, while allowing for specific implementations of regularization techniques.

The optimization algorithms are implemented as separate objects that can be passed to the fit function of the regression classes, promoting modularity and flexibility. The fit function of the model class calls the step function of the optimizer object, which return updated model parameters based on the given optimization algorithm, passing the gradient and the current parameters for each iteration.

```
1 import numpy as np
2 from sklearn.model_selection import train_test_split, KFold
3 from sklearn.utils import resample
4 from .optimization import Optimizer
5
6 class OLS:
7     """
8     Ordinary Least Squares Regression Model.
9     """
10    def __init__(self, x: np.ndarray, y: np.ndarray,
11                 degree: int = 1,
12                 test_size: float = 0.2) -> None:
13        """
14        Initialize the OLS model with data.
15        Args:
16            x (np.ndarray): Input data of shape (n_samples,
17            n_features).
18            y (np.ndarray): Output data of shape (n_samples,)
19            .
20            degree (int): The polynomial degree for the model
21            .
22        """
23        self.x_train, self.x_test, self.y_train, self.y_test
24        = \
25            train_test_split(x, y, test_size=test_size)
26        self.degree = degree
27        self.theta: np.ndarray = None
```

```

24
25     def name(self) -> str:
26         return "OLS"
27
28     def set_degree(self, degree: int) -> None:
29         """
30         Set the polynomial degree for the model.
31         Args:
32             degree (int): The polynomial degree for the model
33         """
34         self.degree = degree
35
36     def design_matrix(self, x: np.ndarray | None = None) ->
37     np.ndarray:
38         """
39         Create the design matrix for polynomial features.
40         Returns:
41             np.ndarray: The design matrix.
42         """
43         if x is None:
44             x = self.x_train
45         return np.vander(x, N=self.degree + 1, increasing=
46         True)
47
48     def fit(self, optimizer: Optimizer | None = None,
49     batch_size: int | None = None) -> None:
50         """
51         Fit the model to the data.
52         Args:
53             optimizer (Optimizer | None): The optimization
54             method to use. If None, use analytical solution.
55             batch_size (int | None): The batch size for
56             stochastic gradient descent. If None, use full batch
57             gradient descent.
58         """
59         if optimizer is None:
60             X = self.design_matrix()
61             self.theta = np.linalg.inv(X.T @ X) @ X.T @ self.
62             y_train # Analytical solution
63         else:
64             self.gradient_descent(optimizer=optimizer,
65             batch_size=batch_size)
66
67     def gradient(self, X: np.ndarray) -> np.ndarray:
68         """
69         Compute the gradient of the loss function.
70         Args:

```

```

64         X (np.ndarray): The design matrix.
65     Returns:
66         np.ndarray: The gradient vector.
67     """
68     return -2/len(self.y_train) * X.T @ (self.y_train - X
@ self.theta)
69
70 def hessian(self, X: np.ndarray) -> np.ndarray:
71     """
72     Compute the Hessian matrix of the loss function.
73     Args:
74         X (np.ndarray): The design matrix.
75     Returns:
76         np.ndarray: The Hessian matrix.
77     """
78     return 2/len(self.y_train) * X.T @ X
79
80 def update_parameters(self, theta: np.ndarray, eta: float
) -> None:
81     """
82     Update the model parameters.
83     Args:
84         theta (np.ndarray): The new parameter values.
85     """
86     self.theta = theta
87
88 def gradient_descent(self, optimizer: Optimizer,
batch_size: int | None = None) -> None:
89     """
90     Fit the model using Gradient Descent.
91     Args:
92         optimizer (Optimizer): The optimization method to
use.
93         batch_size (int | None): The batch size for
stochastic gradient descent. If None, use full batch
gradient descent.
94     """
95     X = self.design_matrix()
96     self.theta = np.zeros_like(X[0]) # Initialize
parameters
97
98     H = self.hessian(X) # Hessian matrix
99     eig = np.linalg.eigvals(H) # Eigenvalues of the
Hessian
100     eta = 1.0 / np.max(eig) # Learning rate based on
Hessian's largest eigenvalue
101
102     optimizer.reset() # Reset optimizer state
for new parameter size

```

```

103         while optimizer.iterate():
104             # Stochastic Gradient Descent with mini-batches
105             if batch_size is not None:
106                 indices = np.random.choice(len(self.y_train),
batch_size, replace=False)
107                 X_batch = X[indices]
108                 y_batch = self.y_train[indices]
109                 grad = -2/batch_size * X_batch.T @ (y_batch -
X_batch @ self.theta)
110             # Full Batch Gradient Descent
111             else:
112                 grad = self.gradient(X)
113
114             # Update parameters using the optimizer
115             step = optimizer.step(self.theta, grad, eta)
116             self.update_parameters(step, eta)
117
118     def predict(self) -> np.ndarray:
119         """
120         Predict using the fitted model.
121         Returns:
122             np.ndarray: Predicted values of shape (n_samples
,).
123         """
124         if self.theta is None:
125             raise ValueError("Model is not fitted yet. Call '
fit' before 'predict'.")
126         X = self.design_matrix(self.x_test)
127         return X @ self.theta
128
129     def mse(self) -> float:
130         """
131         Calculate the Mean Squared Error of the model.
132         Returns:
133             float: Mean Squared Error.
134         """
135         return np.mean((self.y_test - self.predict()) ** 2)
136
137     def r2_score(self) -> float:
138         """
139         Calculate the R score of the model.
140         Returns:
141             float: R score.
142         """
143         return 1 - np.sum((self.y_test - self.predict()) **
2) \
144             / np.sum((self.y_test - np.mean(self.y_test)
) ** 2)
145

```

```

146     def bootstrap(self, optimizer: Optimizer | None = None,
147 n_bootstraps: int = 100) -> tuple[float, float, float]:
148         """
149         Perform bootstrap resampling to estimate bias and
150 variance.
151         Args:
152             n_bootstraps (int): Number of bootstrap samples.
153         Returns:
154             tuple[np.ndarray, np.ndarray, np.ndarray]: Tuple
155 containing
156             - error (np.ndarray): Total error of the
157 model.
158             - bias (np.ndarray): Bias of the model.
159             - variance (np.ndarray): Variance of the
160 model.
161         """
162         x, y = self.x_train, self.y_train # Save original
163 data
164         y_test = self.y_test.reshape(-1, 1)
165         y_pred = np.zeros((self.y_test.shape[0], n_bootstraps
166 ))
167
168         for i in range(n_bootstraps):
169             self.x_train, self.y_train = resample(x, y) #
170 Bootstrap resample
171             self.fit(optimizer=optimizer) # Fit
172 model
173             y_pred[:, i] = self.predict().ravel()
174
175             # Calculate error, bias, and variance
176             error: float = np.mean( (y_test - y_pred)**2 )
177             bias: float = np.mean( (y_test - np.mean(y_pred, axis
178 =1, keepdims=True))**2 )
179             variance: float = np.mean( np.var(y_pred, axis=1,
180 keepdims=True))
181
182             self.x_train, self.y_train = x, y # Restore original
183 data
184
185         return error, bias, variance
186
187     def kfold_cross_validation(self, optimizer: Optimizer |
188 None = None, k: int = 5) -> float:
189         """
190         Perform k-fold cross-validation to estimate the model
191 's performance.
192         Args:
193             k (int): Number of folds.
194         Returns:

```



```

181         float: Average Mean Squared Error across all
folds.
182     """
183
184     kf = KFold(n_splits=k, shuffle=True)
185     mse_list = []
186
187     x, y = self.x_train, self.y_train # Use training
data for cross-validation
188
189     for train_index, val_index in kf.split(x):
190         self.x_train, x_val = x[train_index], x[val_index
]
191         self.y_train, y_val = y[train_index], y[val_index
]
192
193         self.fit(optimizer=optimizer)
194
195         X_val = self.design_matrix(x_val)
196         y_val_pred = X_val @ self.theta
197         mse_fold = np.mean((y_val - y_val_pred) ** 2)
198         mse_list.append(mse_fold)
199
200     self.x_train, self.y_train = x, y # Restore original
training data
201
202     return np.mean(mse_list)

```

Listing 9: models/ols.py

```

1 from typing import override
2 import numpy as np
3
4 from .optimization import Optimizer
5 from .ols import OLS
6
7 class Ridge(OLS):
8     def __init__(self,
9                 x: np.ndarray,
10                y: np.ndarray,
11                degree: int = 1,
12                reg_lambda: float = 1.0,
13                test_size: float = 0.2) -> None:
14
15         """
16         Initialize the Ridge regression model with data.
17         Args:
18             x (np.ndarray): Input data of shape (n_samples,
n_features).
19             y (np.ndarray): Output data of shape (n_samples,)

```

```

19         degree (int): The polynomial degree for the model
20
21         reg_lambda (float): Regularization strength.
22         test_size (float): Proportion of the dataset to
23         include in the test split.
24         """
25         super().__init__(x, y, degree, test_size)
26         self.reg_lambda = reg_lambda
27
28     @override
29     def name(self):
30         return "Ridge"
31
32     def set_lambda(self, reg_lambda: float) -> None:
33         """
34         Set the regularization strength for the model.
35         Args:
36             reg_lambda (float): Regularization strength.
37         """
38         self.reg_lambda = reg_lambda
39
40     @override
41     def fit(self, optimizer: Optimizer | None = None,
42             batch_size: int | None = None) -> None:
43         """
44         Fit the model to the data (analytical).
45         """
46         if optimizer is not None:
47             self.gradient_descent(optimizer=optimizer,
48                                   batch_size=batch_size)
49         else:
50             X = self.design_matrix()
51             self.theta = np.linalg.inv(X.T @ X + self.
52                                         reg_lambda * np.eye(X.shape[1])) @ X.T @ self.y_train
53
54     @override
55     def gradient(self, X):
56         return super().gradient(X) + 2 * self.reg_lambda *
57         self.theta

```

Listing 10: models/ridge.py

```

1 from typing import override
2 import numpy as np
3 from .ols import OLS
4 from .optimization import Optimizer, FixedLearningRate #
5     default optimizer if none provided

```

```

6 class Lasso(OLS):
7     """
8     LASSO (Least Absolute Shrinkage and Selection Operator)
9     Regression Model.
10    """
11    def __init__(self,
12                  x: np.ndarray,
13                  y: np.ndarray,
14                  degree: int = 1,
15                  reg_lambda: float = 0.01,
16                  test_size: float = 0.2) -> None:
17        """
18        Initialize the LASSO regression model with data.
19        Args:
20            x (np.ndarray): Input data of shape (n_samples,
21            n_features).
22            y (np.ndarray): Output data of shape (n_samples,)
23            .
24            degree (int): The polynomial degree for the model
25            .
26            reg_lambda (float): L1 regularization strength (
27            lambda).
28            test_size (float): Proportion of the dataset to
29            include in the test split.
30        """
31        super().__init__(x, y, degree, test_size)
32        self.reg_lambda = reg_lambda
33
34    @override
35    def name(self):
36        return "LASSO"
37
38    def set_lambda(self, reg_lambda: float) -> None:
39        """
40        Set the regularization strength for the model.
41        Args:
42            reg_lambda (float): L1 regularization strength.
43        """
44        self.reg_lambda = reg_lambda
45
46    @override
47    def fit(self, optimizer: Optimizer | None = None,
48            batch_size: int | None = None) -> None:
49        """
50        Fit the LASSO model using gradient descent.
51        LASSO cannot be solved analytically, so it requires
52        an optimizer.
53        Args:

```

```

47         optimizer (Optimizer | None): The optimization
method to use. Required for LASSO.
48         batch_size (int | None): The batch size for
stochastic gradient descent.
49         """
50         if optimizer is None:
51             optimizer = FixedLearningRate()
52
53         self.gradient_descent(optimizer=optimizer, batch_size
=batch_size)
54
55     @override
56     def update_parameters(self, theta: np.ndarray, eta: float
) -> None:
57         """
58         Update the model parameters.
59         Args:
60             theta (np.ndarray): The new parameter values.
61         """
62         self.theta = np.sign(theta) * np.maximum(0, np.abs(
theta) - self.reg_lambda * eta)

```

Listing 11: models/lasso.py

```

1  from typing import override
2  import numpy as np
3
4  MAX_ITER_DEFAULT: int = 1000
5
6  class Optimizer:
7      """
8      Base class for optimization methods.
9      """
10
11     def __init__(self, max_iter: int = MAX_ITER_DEFAULT):
12         self.max_iter = max_iter
13         self.iterations = max_iter
14
15     def iterate(self) -> bool:
16         """
17         Check if more iterations are allowed.
18         Returns:
19             bool: True if more iterations are allowed, False
otherwise.
20         """
21         self.iterations -= 1
22         return self.iterations > 0
23
24     def name(self) -> str:

```

```

25         return self.__class__.__name__
26
27     def step(self, theta: np.ndarray, gradient: np.ndarray,
28 eta: float = 1.0) -> np.ndarray:
29         """
30         Perform a single optimization step.
31         Args:
32             theta (np.ndarray): Current parameters.
33             gradient (np.ndarray): Current gradient.
34             eta (float): Learning rate.
35         Returns:
36             np.ndarray: Updated parameters.
37         """
38         raise NotImplementedError("This method should be
39 overridden by subclasses.")
40
41     def reset(self) -> None:
42         """Reset optimizer state. Override in stateful
43 optimizers."""
44         self.iterations = self.max_iter
45
46 class FixedLearningRate(Optimizer):
47     """
48     Fixed Learning Rate optimization method.
49     """
50
51     @override
52     def name(self) -> str:
53         return "Fixed Learning Rate"
54
55     @override
56     def step(self, theta: np.ndarray, gradient: np.ndarray,
57 eta: float = 1.0) -> np.ndarray:
58         return theta - eta * gradient
59
60 class AdaGrad(Optimizer):
61     """
62     AdaGrad optimization method.
63     """
64
65     def __init__(self, max_iter: int = MAX_ITER_DEFAULT,
66 epsilon: float = 1e-8) -> None:
67         super().__init__(max_iter)
68         self.epsilon = epsilon
69         self.G = None
70
71     @override
72     def name(self) -> str:
73         return "AdaGrad"

```

```

69
70     @override
71     def reset(self) -> None:
72         super().reset()
73         self.G = None
74
75     @override
76     def step(self, theta: np.ndarray, gradient: np.ndarray,
77 eta: float = 1.0) -> np.ndarray:
78         if self.G is None:
79             self.G = np.zeros_like(gradient)
80             self.G += gradient**2
81             lr = eta / (np.sqrt(self.G) + self.epsilon)
82             return theta - lr * gradient
83
84 class Adam(Optimizer):
85     """
86     Adam optimization method.
87     """
88
89     def __init__(self, max_iter: int = MAX_ITER_DEFAULT,
90 beta1: float = 0.9, beta2: float = 0.999, epsilon: float =
91 1e-8) -> None:
92         super().__init__(max_iter)
93         self.beta1 = beta1
94         self.beta2 = beta2
95         self.epsilon = epsilon
96         self.m = None
97         self.v = None
98         self.t = 0
99
100
101     @override
102     def name(self) -> str:
103         return "Adam"
104
105
106     @override
107     def reset(self) -> None:
108         super().reset()
109         self.m = None
110         self.v = None
111         self.t = 0
112
113     @override
114     def step(self, theta: np.ndarray, gradient: np.ndarray,
115 eta: float = 1.0) -> np.ndarray:
116         if self.m is None:
117             self.m = np.zeros_like(gradient)
118         if self.v is None:
119             self.v = np.zeros_like(gradient)

```

```

114         self.t += 1
115         self.m = self.beta1 * self.m + (1 - self.beta1) *
116         gradient
117         self.v = self.beta2 * self.v + (1 - self.beta2) * (
118         gradient ** 2)
119
120         m_hat = self.m / (1 - self.beta1 ** self.t)
121         v_hat = self.v / (1 - self.beta2 ** self.t)
122
123         return theta - (eta * m_hat) / (np.sqrt(v_hat) + self
124         .epsilon)
125
126 class Momentum(Optimizer):
127     """
128     Momentum optimization method.
129     """
130
131     def __init__(self, max_iter: int = MAX_ITER_DEFAULT,
132     momentum: float = 0.9) -> None:
133         super().__init__(max_iter)
134         self.momentum = momentum
135         self.velocity = None
136
137     @override
138     def name(self) -> str:
139         return "Momentum"
140
141     @override
142     def reset(self) -> None:
143         super().reset()
144         self.velocity = None
145
146     @override
147     def step(self, theta: np.ndarray, gradient: np.ndarray,
148     eta: float = 1.0) -> np.ndarray:
149         if self.velocity is None:
150             self.velocity = np.zeros_like(gradient)
151         self.velocity = self.momentum * self.velocity - eta *
152         gradient
153         return theta + self.velocity
154
155 class RMSProp(Optimizer):
156     """
157     RMSProp optimization method.
158     """
159
160     def __init__(self, max_iter: int = MAX_ITER_DEFAULT,
161     decay_rate: float = 0.9, epsilon: float = 1e-8) -> None:

```

```

156         super().__init__(max_iter)
157         self.decay_rate = decay_rate
158         self.epsilon = epsilon
159         self.cache = None
160
161     @override
162     def name(self) -> str:
163         return "RMSProp"
164
165     @override
166     def reset(self) -> None:
167         super().reset()
168         self.cache = None
169
170     @override
171     def step(self, theta: np.ndarray, gradient: np.ndarray,
172             eta: float = 1.0) -> np.ndarray:
173         if self.cache is None:
174             self.cache = np.zeros_like(gradient)
175             self.cache = self.decay_rate * self.cache + (1 - self
176             .decay_rate) * (gradient ** 2)
177             lr = eta / (np.sqrt(self.cache) + self.epsilon)
178             return theta - lr * gradient

```

Listing 12: models/optimizer.py

```

1 import numpy as np
2
3 START = -1.0
4 STOP = 1.0
5
6 def runge(x: np.ndarray, noise: bool = False) -> np.ndarray:
7     """
8     Compute the Runge function with optional Gaussian noise.
9     Args:
10         x (np.ndarray): Input data of shape (n_samples,).
11         noise (bool): If True, add Gaussian noise to the
12         output.
13
14     Returns:
15         np.ndarray: Output data of shape (n_samples,).
16     """
17     y = 1 / (1 + 25 * x**2)
18     if noise:
19         y += 0.1 * np.random.normal(0, 1, x.shape)
20     return y
21
22 def x(n_samples: int = 100, random_state: int | None = None)
23     -> np.ndarray:

```



```

22     """
23     Generate uniformly spaced input data in the range [start,
24     stop].
25     Args:
26         start (float): Start of the range.
27         stop (float): End of the range.
28         n_samples (int): Number of samples to generate.
29         random_state (int | None): Seed for the random number
30         generator.
31     Returns:
32         np.ndarray: Input data of shape (n_samples,).
33     """
34     if random_state is not None:
35         np.random.seed(random_state)
36     return np.linspace(START, STOP, n_samples)
37
38 def generate_data(n_samples: int = 1000, noise: bool = True,
39 random_state: int | None = None) -> tuple[np.ndarray, np.
40 ndarray]:
41     """
42     Generate input and output data using the Runge function.
43     Args:
44         n_samples (int): Number of samples to generate.
45         noise (bool): If True, add Gaussian noise to the
46         output.
47         random_state (int | None): Seed for the random number
48         generator.
49     Returns:
50         tuple[np.ndarray, np.ndarray]: Tuple containing input
51         data (x) and output data (y).
52     """
53     x_data = x(n_samples=n_samples, random_state=random_state)
54     y_data = runge(x_data, noise=noise)
55     return x_data, y_data

```

Listing 13: data.py

```

1 import numpy as np
2 from matplotlib import pyplot as plt
3
4 from data import generate_data
5 from models.ols import OLS
6 from models.ridge import Ridge
7 from models.lasso import Lasso
8 from models.optimization import (
9     Optimizer,

```

```

10     FixedLearningRate,
11     AdaGrad,
12     Adam,
13     RMSProp
14 )
15
16 def visualize_data(x: np.ndarray, y: np.ndarray,
17                  model: OLS,
18                  optimizer: Optimizer | None = None,
19                  batch_size: int | None = None) -> None:
20     """
21     Visualize the input data and a simple OLS fit.
22     Args:
23         x (np.ndarray): Input data of shape (n_samples,).
24         y (np.ndarray): Output data of shape (n_samples,).
25         model (OLS): The OLS model to fit and visualize.
26         optimizer (Optimizer | None): The optimization method
27         to use. If None, use analytical solution.
28         batch_size (int | None): The batch size for
29         stochastic gradient descent. If None,
30         Returns:
31         None
32     """
33     name = model.name()
34     if optimizer is None:
35         name += "(Analytical Solution)"
36     else:
37         name += f"(Gradient Descent with {optimizer.name()})"
38         if batch_size is not None:
39             name += f", batch_size={batch_size}"
40
41     # fit model and predict
42     model.fit(optimizer=optimizer, batch_size=batch_size)
43     y_pred = model.predict()
44
45     # Create figure or use existing one
46     plt.figure(figsize=(10, 6))
47     # Plot data and prediction points
48     plt.scatter(x, y, alpha=0.6, color='lightgray', s=20,
49                label='Data', zorder=1)
50     plt.scatter(model.x_test, y_pred, label='Prediction',
51                color='blue', zorder=2)
52
53     # Create smooth prediction line
54     x_smooth = np.linspace(x.min(), x.max(), 300)
55     X_smooth = np.vander(x_smooth, N=model.degree + 1,
56                          increasing=True)
57     y_smooth = X_smooth @ model.theta
58     plt.plot(x_smooth, y_smooth, color='red', linewidth=2,

```

```

label='Prediction', zorder=2)

54
55 # Final plot settings
56 plt.title(f"{name}")
57 plt.xlabel("x")
58 plt.ylabel("y")
59 plt.legend()
60 plt.grid()
61 plt.tight_layout()
62 plt.show()
63
64 def analyze_bias_variance(x: np.ndarray, y: np.ndarray,
65                           max_degree: int = 15,
66                           n_bootstraps: int = 100) -> None:
67     """
68     Analyze bias-variance trade-off using bootstrap
69     resampling.
70     Args:
71         x (np.ndarray): Input data of shape (n_samples,).
72         y (np.ndarray): Output data of shape (n_samples,).
73         degree (int): Polynomial degree for the model.
74         n_bootstraps (int): Number of bootstrap samples.
75     """
76     from models.ols import OLS
77     model: OLS = OLS(x, y)
78
79     degrees: np.ndarray = np.arange(1, max_degree + 1)
80     error: np.ndarray = np.zeros(max_degree)
81     bias: np.ndarray = np.zeros(max_degree)
82     variance: np.ndarray = np.zeros(max_degree)
83
84     for degree in degrees:
85         model.set_degree(degree)
86         error[degree - 1], bias[degree - 1], variance[degree
87 - 1] = \
88             model.bootstrap(n_bootstraps=n_bootstraps)
89
90     plt.figure(figsize=(8, 6))
91     plt.plot(degrees, error, marker='o', label='Total Error')
92     plt.plot(degrees, bias, marker='o', label='Bias ')
93     plt.plot(degrees, variance, marker='o', label='Variance')
94     plt.title("Bias-Variance Trade-off")
95     plt.xlabel("Degree")
96     plt.ylabel("Error")
97     plt.legend()
98     plt.grid()
99     plt.show()
100
101 def analyze_cross_validation(x: np.ndarray, y: np.ndarray,

```

```

100         models: list[OLS],
101         max_degree: int = 15,
102         k_folds: int = 5,
103         optimizer: Optimizer | None =
None) -> None:
104     """
105     Analyze model performance using k-fold cross-validation.
106     Args:
107         x (np.ndarray): Input data of shape (n_samples,).
108         y (np.ndarray): Output data of shape (n_samples,).
109         models (list[OLS]): List of models to evaluate.
110         max_degree (int): Maximum polynomial degree to
evaluate.
111         k_folds (int): Number of folds for cross-validation.
112     """
113
114     degrees: np.ndarray = np.arange(1, max_degree + 1)
115     mse_cv: np.ndarray = np.zeros((len(models), max_degree))
116
117     for i, model in enumerate(models):
118         for degree in degrees:
119             model.set_degree(degree)
120             mse_cv[i, degree - 1] = model.
kfold_cross_validation(optimizer=optimizer, k=k_folds)
121
122     plt.figure(figsize=(10, 6))
123     for i, model in enumerate(models):
124         plt.plot(degrees, mse_cv[i], marker='o', label=f"{
model.name()} = {getattr(model, 'reg_lambda', 'N/A')}")
125     plt.title(f"Mean Squared Error - Cross Validation ({
k_folds}-fold), Optimizer: {optimizer.name() if optimizer
else 'Analytical'}")
126     plt.xlabel("Degree")
127     plt.ylabel("MSE")
128     plt.grid()
129     plt.legend()
130     plt.show()
131
132 if __name__ == "__main__":
133     # Generate data
134     n = 100
135     x, y = generate_data(n_samples=n, noise=True,
random_state=42)
136     optimizer = FixedLearningRate()
137
138     # Visualize data with OLS fit
139     ols = OLS(x, y, degree=8)
140     visualize_data(x, y, ols, optimizer=optimizer)
141

```

```

142     # Visualize data with Ridge fit
143     ridge = Ridge(x, y, degree=8, reg_lambda=0.0001)
144     visualize_data(x, y, ridge, optimizer=optimizer)
145
146     # Visualize data with Lasso fit
147     lasso = Lasso(x, y, degree=8, reg_lambda=0.001)
148     visualize_data(x, y, lasso, optimizer=optimizer)
149
150     # Cross-Validation
151     models: list[OLS] = [ols, ridge, lasso]
152     analyze_cross_validation(x, y, models=models, max_degree
=20, k_folds=5, optimizer=optimizer)
153
154     # Analyze bias-variance trade-off
155     analyze_bias_variance(x, y, max_degree=16, n_bootstraps
=100)

```

Listing 14: main.py